

# Pretrained\_VS\_CNN

January 20, 2025

Difference Between CNNs and Pre-trained Models

1. CNNs (Convolutional Neural Networks):  
Definition: A type of neural network specifically designed to process grid-like data such as images. It is built from scratch using components like convolutional layers, pooling layers, and fully connected layers for classification. Advantages: Highly flexible for designing custom architectures tailored to specific datasets. Useful for understanding the impact of architectural modifications on model performance. Challenges: Requires training from scratch, which demands large datasets and significant computational resources.
2. Pre-trained Models: Definition: Models that have been trained on large datasets (e.g., ImageNet) and come with pre-learned weights that can be reused for similar tasks. Advantages: Reduces the need for large amounts of data. Provides a quick way to achieve high performance without training a model from scratch. Challenges: May not perform optimally if the new task is very different from the original task the model was trained on.

Key Difference:

Custom CNNs: Built and trained from scratch, offering complete flexibility but requiring substantial resources.

Pre-trained Models: Start with pre-learned weights, saving time and computational effort, and can be fine-tuned for the specific task.

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

So lets try to read some data and start with CNN

## 1 Split data in drive into Train,Test and Validation

```
[ ]: import os
import shutil
from sklearn.model_selection import train_test_split

def split_data(data_dir, output_dir, train_ratio=0.7, val_ratio=0.25,
               test_ratio=0.05):
    """
    Split data into train, validation, and test directories.

    Parameters:
        data_dir (str): Path to the dataset directory.
```

```

        output_dir (str): Path to save the split dataset.
        train_ratio (float): Ratio of training data (default is 0.7).
        val_ratio (float): Ratio of validation data (default is 0.2).
        test_ratio (float): Ratio of test data (default is 0.1).
    """
    if not os.path.exists(output_dir):
        os.makedirs(output_dir)

    classes = os.listdir(data_dir)
    for class_name in classes:
        class_dir = os.path.join(data_dir, class_name)
        if not os.path.isdir(class_dir):
            continue

        # Get all files in the class directory
        files = os.listdir(class_dir)
        train_files, temp_files = train_test_split(files, test_size=(1 -
↪train_ratio), random_state=42)
        val_files, test_files = train_test_split(temp_files,
↪test_size=(test_ratio / (val_ratio + test_ratio)), random_state=42)

        # Define output subdirectories
        train_dir = os.path.join(output_dir, 'train', class_name)
        val_dir = os.path.join(output_dir, 'val', class_name)
        test_dir = os.path.join(output_dir, 'test', class_name)
        os.makedirs(train_dir, exist_ok=True)
        os.makedirs(val_dir, exist_ok=True)
        os.makedirs(test_dir, exist_ok=True)

        # Move files to respective directories
        for file in train_files:
            shutil.copy(os.path.join(class_dir, file), os.path.join(train_dir,
↪file))
        for file in val_files:
            shutil.copy(os.path.join(class_dir, file), os.path.join(val_dir,
↪file))
        for file in test_files:
            shutil.copy(os.path.join(class_dir, file), os.path.join(test_dir,
↪file))

    print("Data split completed!")

```

```

[ ]: '''
data_dir = '/content/drive/MyDrive/Flowers/train' # Path to your dataset
output_dir = '/content/drive/MyDrive/Flowers_Divided' # Path to save split data
split_data(data_dir, output_dir)'''

```

```
[ ]: "\ndata_dir = '/content/drive/MyDrive/Flowers/train' # Path to your
dataset\noutput_dir = '/content/drive/MyDrive/Flowers_Divided' # Path to save
split data\nsplit_data(data_dir, output_dir)"
```

```
[ ]: import os

def dataset_info(data_dir):
    """
    Display information about the dataset: folder structure and number of
    ↪ images per class.
    This function will now be applied to the train, validation, and test
    ↪ directories.

    Parameters:
        data_dir (str): Path to the dataset directory (train, val, or test).
    """
    total_images = 0
    class_counts = {}

    # Traverse dataset directory (train, val, or test)
    for class_name in sorted(os.listdir(data_dir)):
        class_dir = os.path.join(data_dir, class_name)
        if not os.path.isdir(class_dir):
            continue # Skip non-directory files

        # Count images in the current class folder
        num_images = len([f for f in os.listdir(class_dir) if os.path.isfile(os.
    ↪ path.join(class_dir, f))])
        class_counts[class_name] = num_images
        total_images += num_images

    # Print results
    print(f"Dataset Directory: {data_dir}")
    print(f"Total Classes: {len(class_counts)}")
    print(f"Total Images: {total_images}\n")

    print("Images per Class:")
    for class_name, count in class_counts.items():
        print(f" {class_name}: {count} images")

# Example usage
train_dir = '/content/drive/MyDrive/divided_mask_dataset/train' # Path to your
    ↪ train data
val_dir = '/content/drive/My Drive/divided_mask_dataset/val' # Path to
    ↪ your validation data
test_dir = '/content/drive/My Drive/divided_mask_dataset/test' # Path to
    ↪ your test data
```

```

print("Train Dataset Info:")
dataset_info(train_dir)

print("\nValidation Dataset Info:")
dataset_info(val_dir)

print("\nTest Dataset Info:")
dataset_info(test_dir)

```

Train Dataset Info:  
Dataset Directory: /content/drive/MyDrive/divided\_mask\_dataset/train  
Total Classes: 2  
Total Images: 962

Images per Class:  
    with mask: 482 images  
    without mask: 480 images

Validation Dataset Info:  
Dataset Directory: /content/drive/My Drive/divided\_mask\_dataset/val  
Total Classes: 2  
Total Images: 275

Images per Class:  
    with mask: 138 images  
    without mask: 137 images

Test Dataset Info:  
Dataset Directory: /content/drive/My Drive/divided\_mask\_dataset/test  
Total Classes: 2  
Total Images: 139

Images per Class:  
    with mask: 70 images  
    without mask: 69 images

plot some images

```

[ ]: import os
import random
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

def plot_random_images(data_dir, num_images=5):
    """
    Plot random images from the specified dataset directory.

```

```

Parameters:
    data_dir (str): Path to the dataset directory (train, val, or test).
    num_images (int): Number of random images to display.
"""
class_names = sorted(os.listdir(data_dir))
class_name = random.choice(class_names) # Choose a random class to display
↪ images
class_dir = os.path.join(data_dir, class_name)

# Get all image file names in the chosen class folder
image_files = [f for f in os.listdir(class_dir) if os.path.isfile(os.path.
↪ join(class_dir, f))]

# Choose random images to display
selected_images = random.sample(image_files, num_images)

# Plot the selected images
plt.figure(figsize=(15, 10))
for i, img_file in enumerate(selected_images):
    img_path = os.path.join(class_dir, img_file)
    img = mpimg.imread(img_path)

    plt.subplot(1, num_images, i+1)
    plt.imshow(img)
    plt.axis('off') # Hide axis
    plt.title(f"Class: {class_name}\nImage: {img_file}")

plt.show()

print("Random Images from Train Dataset:")
plot_random_images(train_dir)

print("Random Images from Validation Dataset:")
plot_random_images(val_dir)

print("Random Images from Test Dataset:")
plot_random_images(test_dir)

```

Random Images from Train Dataset:



Random Images from Validation Dataset:



Random Images from Test Dataset:



Create ImageDataGenerators for train, validation, and test sets note : Val and test have no augmentation

```
[ ]: import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Define the target size for the CNN model
```

```

target_size = (224, 224) # Target size for models like MobileNetV2

# Define the batch size
batch_size = 32

# Create ImageDataGenerators for train, validation, and test sets
train_datagen = ImageDataGenerator(
    rescale=1./255, # Rescale pixel values to [0, 1]
    rotation_range=20, # Optional: Data augmentation
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

val_test_datagen = ImageDataGenerator(rescale=1./255) # No augmentation for
↳ validation/test

# Create generators to load and resize images
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=target_size, # Resize images to 224x224
    batch_size=batch_size,
    class_mode='categorical', # For multi-class classification
    shuffle=True
)

val_generator = val_test_datagen.flow_from_directory(
    val_dir,
    target_size=target_size, # Resize images to 224x224
    batch_size=batch_size,
    class_mode='categorical',
    shuffle=False
)

test_generator = val_test_datagen.flow_from_directory(
    test_dir,
    target_size=target_size, # Resize images to 224x224
    batch_size=batch_size,
    class_mode='categorical',
    shuffle=False
)

```

Found 962 images belonging to 2 classes.  
Found 275 images belonging to 2 classes.  
Found 139 images belonging to 2 classes.

## 2 Build the CNN

```
[ ]: # Build a simple CNN model
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(2, activation='softmax') # 2 classes for CIFAR-10
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

[ ]: from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

early_stopping = EarlyStopping(monitor='val_loss', patience=5,
                               restore_best_weights=True)
model_checkpoint = ModelCheckpoint('/content/drive/MyDrive/divided_mask_dataset/
                               best_advanced_model.keras', monitor='val_loss', save_best_only=True,
                               verbose=1)

[ ]: # Train the model using the generators
history = model.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // batch_size,
    epochs=50,
    validation_data=val_generator,
    validation_steps=val_generator.samples // batch_size,
    callbacks=[early_stopping, model_checkpoint]
)
```

Epoch 1/50

/usr/local/lib/python3.10/dist-

packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:122:

UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self.\_warn\_if\_super\_not\_called()

30/30 0s 476ms/step -

accuracy: 0.5332 - loss: 3.6778

Epoch 1: val\_loss improved from inf to 0.64849, saving model to



```

/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
30/30          30s 638ms/step -
accuracy: 0.5326 - loss: 3.6211 - val_accuracy: 0.5469 - val_loss: 0.6485
Epoch 2/50
  1/30          1s 52ms/step - accuracy:
0.5312 - loss: 0.6448

/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data;
interrupting training. Make sure that your dataset or generator can generate at
least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()`
function when building your dataset.
  self.gen.throw(typ, value, traceback)

Epoch 2: val_loss did not improve from 0.64849
30/30          2s 57ms/step -
accuracy: 0.5312 - loss: 0.6448 - val_accuracy: 0.0000e+00 - val_loss: 0.9708
Epoch 3/50
30/30          0s 435ms/step -
accuracy: 0.7363 - loss: 0.5029
Epoch 3: val_loss improved from 0.64849 to 0.40754, saving model to
/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
30/30          20s 545ms/step -
accuracy: 0.7380 - loss: 0.5018 - val_accuracy: 0.8125 - val_loss: 0.4075
Epoch 4/50
  1/30          1s 52ms/step - accuracy:
0.8438 - loss: 0.2785
Epoch 4: val_loss did not improve from 0.40754
30/30          2s 64ms/step -
accuracy: 0.8438 - loss: 0.2785 - val_accuracy: 0.7368 - val_loss: 0.7172
Epoch 5/50
29/30          0s 412ms/step -
accuracy: 0.8509 - loss: 0.3546
Epoch 5: val_loss improved from 0.40754 to 0.17103, saving model to
/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
30/30          37s 513ms/step -
accuracy: 0.8517 - loss: 0.3541 - val_accuracy: 0.9492 - val_loss: 0.1710
Epoch 6/50
  1/30          1s 52ms/step - accuracy:
0.7812 - loss: 0.4487
Epoch 6: val_loss did not improve from 0.17103
30/30          2s 60ms/step -
accuracy: 0.7812 - loss: 0.4487 - val_accuracy: 0.8947 - val_loss: 0.3407
Epoch 7/50
30/30          0s 456ms/step -
accuracy: 0.8883 - loss: 0.3253
Epoch 7: val_loss improved from 0.17103 to 0.15580, saving model to
/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
30/30          20s 564ms/step -

```

accuracy: 0.8884 - loss: 0.3248 - val\_accuracy: 0.9688 - val\_loss: 0.1558  
 Epoch 8/50  
   1/30                   0s 13ms/step - accuracy:  
 0.5000 - loss: 0.7200  
 Epoch 8: val\_loss did not improve from 0.15580  
 30/30                   0s 1ms/step -  
 accuracy: 0.5000 - loss: 0.7200 - val\_accuracy: 0.8947 - val\_loss: 0.5462  
 Epoch 9/50  
   29/30                   0s 432ms/step -  
 accuracy: 0.9014 - loss: 0.3270  
 Epoch 9: val\_loss improved from 0.15580 to 0.15452, saving model to  
 /content/drive/MyDrive/divided\_mask\_dataset/best\_advanced\_model.keras  
 30/30                   21s 534ms/step -  
 accuracy: 0.9013 - loss: 0.3277 - val\_accuracy: 0.9570 - val\_loss: 0.1545  
 Epoch 10/50  
   1/30                   1s 52ms/step - accuracy:  
 0.9062 - loss: 0.3866  
 Epoch 10: val\_loss did not improve from 0.15452  
 30/30                   0s 1ms/step -  
 accuracy: 0.9062 - loss: 0.3866 - val\_accuracy: 0.8947 - val\_loss: 0.4882  
 Epoch 11/50  
   29/30                   0s 373ms/step -  
 accuracy: 0.9258 - loss: 0.2594  
 Epoch 11: val\_loss improved from 0.15452 to 0.13850, saving model to  
 /content/drive/MyDrive/divided\_mask\_dataset/best\_advanced\_model.keras  
 30/30                   40s 495ms/step -  
 accuracy: 0.9252 - loss: 0.2604 - val\_accuracy: 0.9609 - val\_loss: 0.1385  
 Epoch 12/50  
   1/30                   1s 51ms/step - accuracy:  
 0.9062 - loss: 0.2503  
 Epoch 12: val\_loss did not improve from 0.13850  
 30/30                   0s 887us/step -  
 accuracy: 0.9062 - loss: 0.2503 - val\_accuracy: 0.8947 - val\_loss: 0.3852  
 Epoch 13/50  
   29/30                   0s 462ms/step -  
 accuracy: 0.9118 - loss: 0.2450  
 Epoch 13: val\_loss improved from 0.13850 to 0.11348, saving model to  
 /content/drive/MyDrive/divided\_mask\_dataset/best\_advanced\_model.keras  
 30/30                   21s 578ms/step -  
 accuracy: 0.9121 - loss: 0.2439 - val\_accuracy: 0.9648 - val\_loss: 0.1135  
 Epoch 14/50  
   1/30                   1s 52ms/step - accuracy:  
 0.9688 - loss: 0.0880  
 Epoch 14: val\_loss did not improve from 0.11348  
 30/30                   1s 43ms/step -  
 accuracy: 0.9688 - loss: 0.0880 - val\_accuracy: 0.8947 - val\_loss: 0.3012  
 Epoch 15/50  
   30/30                   0s 437ms/step -

```

accuracy: 0.9468 - loss: 0.1869
Epoch 15: val_loss did not improve from 0.11348
30/30          18s 492ms/step -
accuracy: 0.9466 - loss: 0.1875 - val_accuracy: 0.9648 - val_loss: 0.1456
Epoch 16/50
 1/30          1s 45ms/step - accuracy:
0.9375 - loss: 0.1846
Epoch 16: val_loss improved from 0.11348 to 0.10635, saving model to
/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
30/30          3s 90ms/step -
accuracy: 0.9375 - loss: 0.1846 - val_accuracy: 1.0000 - val_loss: 0.1063
Epoch 17/50
29/30          0s 455ms/step -
accuracy: 0.9342 - loss: 0.2232
Epoch 17: val_loss improved from 0.10635 to 0.07669, saving model to
/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
30/30          19s 546ms/step -
accuracy: 0.9339 - loss: 0.2233 - val_accuracy: 0.9766 - val_loss: 0.0767
Epoch 18/50
 1/30          1s 53ms/step - accuracy:
0.9688 - loss: 0.0839
Epoch 18: val_loss did not improve from 0.07669
30/30          0s 947us/step -
accuracy: 0.9688 - loss: 0.0839 - val_accuracy: 0.8947 - val_loss: 0.3288
Epoch 19/50
29/30          0s 659ms/step -
accuracy: 0.9308 - loss: 0.2163
Epoch 19: val_loss improved from 0.07669 to 0.06738, saving model to
/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
30/30          26s 761ms/step -
accuracy: 0.9311 - loss: 0.2153 - val_accuracy: 0.9844 - val_loss: 0.0674
Epoch 20/50
 1/30          1s 52ms/step - accuracy:
1.0000 - loss: 0.1187
Epoch 20: val_loss did not improve from 0.06738
30/30          0s 851us/step -
accuracy: 1.0000 - loss: 0.1187 - val_accuracy: 0.8947 - val_loss: 0.2825
Epoch 21/50
29/30          0s 451ms/step -
accuracy: 0.9420 - loss: 0.1719
Epoch 21: val_loss did not improve from 0.06738
30/30          18s 484ms/step -
accuracy: 0.9419 - loss: 0.1729 - val_accuracy: 0.9766 - val_loss: 0.0897
Epoch 22/50
 1/30          1s 40ms/step - accuracy:
1.0000 - loss: 0.0962
Epoch 22: val_loss did not improve from 0.06738
30/30          0s 2ms/step -

```

```

accuracy: 1.0000 - loss: 0.0962 - val_accuracy: 0.8947 - val_loss: 0.1552
Epoch 23/50
30/30          0s 381ms/step -
accuracy: 0.9231 - loss: 0.1930
Epoch 23: val_loss improved from 0.06738 to 0.06149, saving model to
/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
30/30          19s 512ms/step -
accuracy: 0.9228 - loss: 0.1947 - val_accuracy: 0.9805 - val_loss: 0.0615
Epoch 24/50
1/30           1s 57ms/step - accuracy:
0.9688 - loss: 0.1414
Epoch 24: val_loss did not improve from 0.06149
30/30          1s 43ms/step -
accuracy: 0.9688 - loss: 0.1414 - val_accuracy: 0.8947 - val_loss: 0.1520
Epoch 25/50
29/30          0s 453ms/step -
accuracy: 0.9036 - loss: 0.2319
Epoch 25: val_loss did not improve from 0.06149
30/30          18s 486ms/step -
accuracy: 0.9044 - loss: 0.2310 - val_accuracy: 0.9688 - val_loss: 0.0989
Epoch 26/50
1/30           1s 39ms/step - accuracy:
1.0000 - loss: 0.1066
Epoch 26: val_loss did not improve from 0.06149
30/30          1s 39ms/step -
accuracy: 1.0000 - loss: 0.1066 - val_accuracy: 0.8947 - val_loss: 0.3627
Epoch 27/50
29/30          0s 415ms/step -
accuracy: 0.9354 - loss: 0.1990
Epoch 27: val_loss did not improve from 0.06149
30/30          19s 450ms/step -
accuracy: 0.9356 - loss: 0.1984 - val_accuracy: 0.9883 - val_loss: 0.0640
Epoch 28/50
1/30           1s 40ms/step - accuracy:
1.0000 - loss: 0.1027
Epoch 28: val_loss did not improve from 0.06149
30/30          1s 39ms/step -
accuracy: 1.0000 - loss: 0.1027 - val_accuracy: 0.8947 - val_loss: 0.1296

```

## 2.1 Evaluate

```

[ ]: # Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_generator, steps=test_generator.
    ↪samples // batch_size)
print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_accuracy}")

```

```

4/4          1s 246ms/step -

```

accuracy: 0.9885 - loss: 0.0422  
Test Loss: 0.07171785086393356  
Test Accuracy: 0.9765625

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
from tensorflow.keras.models import load_model

# Generate predictions for the test data without specifying 'steps'
test_predictions = model.predict(test_generator, verbose=1)

# Get the predicted class labels
predicted_labels = np.argmax(test_predictions, axis=1)


# Get the true class labels from the test generator
true_labels = test_generator.classes

# Generate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(10, 7))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=test_generator.
    ↪class_indices.keys(), yticklabels=test_generator.class_indices.keys())
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

5/5

4s 778ms/step



i expected complexed data than what i see but lets try  
 another confused data to see how CNN dealing with  
 Confused classes

### 3 Try confused data

Flowers data may be difficult for Simple CNN to Classifing

```
[ ]: import os

def dataset_info(data_dir):
    """
    Display information about the dataset: folder structure and number of
    ↪ images per class.
    This function will now be applied to the train, validation, and test
    ↪ directories.

    Parameters:
```

```

    data_dir (str): Path to the dataset directory (train, val, or test).
    """
    total_images = 0
    class_counts = {}

    # Traverse dataset directory (train, val, or test)
    for class_name in sorted(os.listdir(data_dir)):
        class_dir = os.path.join(data_dir, class_name)
        if not os.path.isdir(class_dir):
            continue # Skip non-directory files

        # Count images in the current class folder
        num_images = len([f for f in os.listdir(class_dir) if os.path.isfile(os.
↪path.join(class_dir, f))])
        class_counts[class_name] = num_images
        total_images += num_images

    # Print results
    print(f"Dataset Directory: {data_dir}")
    print(f"Total Classes: {len(class_counts)}")
    print(f"Total Images: {total_images}\n")

    print("Images per Class:")
    for class_name, count in class_counts.items():
        print(f"  {class_name}: {count} images")

# Example usage
train_dir = '/content/drive/MyDrive/Flowers_Divided/train' # Path to your ↵
↪train data
val_dir = '/content/drive/MyDrive/Flowers_Divided/val' # Path to your ↵
↪validation data
test_dir = '/content/drive/MyDrive/Flowers_Divided/test' # Path to your test ↵
↪data

print("Train Dataset Info:")
dataset_info(train_dir)

print("\nValidation Dataset Info:")
dataset_info(val_dir)

print("\nTest Dataset Info:")
dataset_info(test_dir)

```

```

Train Dataset Info:
Dataset Directory: /content/drive/MyDrive/Flowers_Divided/train
Total Classes: 5
Total Images: 1919

```

Images per Class:

daisy: 350 images  
dandelion: 452 images  
rose: 347 images  
sunflower: 346 images  
tulip: 424 images

Validation Dataset Info:

Dataset Directory: /content/drive/MyDrive/Flowers\_Divided/val  
Total Classes: 5  
Total Images: 686

Images per Class:

daisy: 125 images  
dandelion: 161 images  
rose: 124 images  
sunflower: 124 images  
tulip: 152 images

Test Dataset Info:

Dataset Directory: /content/drive/MyDrive/Flowers\_Divided/test  
Total Classes: 5  
Total Images: 141

Images per Class:

daisy: 26 images  
dandelion: 33 images  
rose: 26 images  
sunflower: 25 images  
tulip: 31 images

```
[ ]: import os
import random
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

def plot_random_images(data_dir, num_images=5):
    """
    Plot random images from the specified dataset directory.

    Parameters:
        data_dir (str): Path to the dataset directory (train, val, or test).
        num_images (int): Number of random images to display.
    """
    class_names = sorted(os.listdir(data_dir))
```



```

class_name = random.choice(class_names) # Choose a random class to display
↪ images
class_dir = os.path.join(data_dir, class_name)

# Get all image file names in the chosen class folder
image_files = [f for f in os.listdir(class_dir) if os.path.isfile(os.path.
↪ join(class_dir, f))]

# Choose random images to display
selected_images = random.sample(image_files, num_images)

# Plot the selected images
plt.figure(figsize=(15, 10))
for i, img_file in enumerate(selected_images):
    img_path = os.path.join(class_dir, img_file)
    img = mpimg.imread(img_path)

    plt.subplot(1, num_images, i+1)
    plt.imshow(img)
    plt.axis('off') # Hide axis
    plt.title(f"Class: {class_name}")

plt.show()

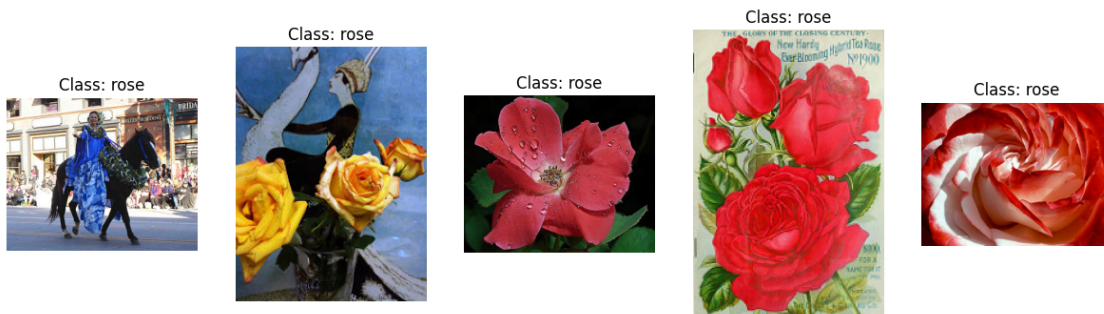
print("Random Images from Train Dataset:")
plot_random_images(train_dir)

print("Random Images from Validation Dataset:")
plot_random_images(val_dir)

print("Random Images from Test Dataset:")
plot_random_images(test_dir)

```


Random Images from Train Dataset:



Random Images from Validation Dataset:



Random Images from Test Dataset:



```
[ ]: import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Define the target size for the CNN model
target_size = (224, 224) # Target size for models like MobileNetV2

# Define the batch size
batch_size = 32

# Create ImageDataGenerators for train, validation, and test sets
train_datagen = ImageDataGenerator(
    rescale=1./255, # Rescale pixel values to [0, 1]
    rotation_range=20, # Optional: Data augmentation
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

val_test_datagen = ImageDataGenerator(rescale=1./255) # No augmentation for
↳ validation/test
```

```

# Create generators to load and resize images
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=target_size, # Resize images to 224x224
    batch_size=batch_size,
    class_mode='categorical', # For multi-class classification (CIFAR-10)
    shuffle=True
)

val_generator = val_test_datagen.flow_from_directory(
    val_dir,
    target_size=target_size, # Resize images to 224x224
    batch_size=batch_size,
    class_mode='categorical',
    shuffle=False
)

test_generator = val_test_datagen.flow_from_directory(
    test_dir,
    target_size=target_size, # Resize images to 224x224
    batch_size=batch_size,
    class_mode='categorical',
    shuffle=False
)

```

Found 1919 images belonging to 5 classes.

Found 686 images belonging to 5 classes.

Found 141 images belonging to 5 classes.

```

[ ]: # Build a simple CNN model
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(5, activation='softmax') # 10 classes for CIFAR-10
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

/usr/local/lib/python3.10/dist-

packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not

pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[ ]: from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

early_stopping = EarlyStopping(monitor='val_loss', patience=5,
    ↳ restore_best_weights=True)
model_checkpoint = ModelCheckpoint('/content/drive/MyDrive/divided_mask_dataset/
    ↳ best_advanced_model.keras', monitor='val_loss', save_best_only=True,
    ↳ verbose=1)
```

```
[ ]: # Train the model using the generators
history = model.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // batch_size,
    epochs=50,
    validation_data=val_generator,
    validation_steps=val_generator.samples // batch_size,
    callbacks=[early_stopping, model_checkpoint]
)
```

Epoch 1/50

/usr/local/lib/python3.10/dist-

packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:122:

UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.

```
self._warn_if_super_not_called()
```

59/59 0s 618ms/step -

accuracy: 0.2674 - loss: 4.8076

Epoch 1: val\_loss improved from inf to 1.26450, saving model to

/content/drive/MyDrive/divided\_mask\_dataset/best\_advanced\_model.keras

59/59 64s 893ms/step -

accuracy: 0.2685 - loss: 4.7680 - val\_accuracy: 0.4628 - val\_loss: 1.2645

Epoch 2/50

1/59 4:07 4s/step - accuracy:

0.5806 - loss: 1.2086

/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps\_per\_epoch \* epochs` batches. You may need to use the `.repeat()` function when building your dataset.

```
self.gen.throw(typ, value, traceback)
```

```

Epoch 2: val_loss improved from 1.26450 to 0.97379, saving model to
/content/drive/MyDrive/divided_mask_dataset/best_advanced_model.keras
59/59          28s 405ms/step -
accuracy: 0.5806 - loss: 1.2086 - val_accuracy: 0.3571 - val_loss: 0.9738
Epoch 3/50
58/59          0s 496ms/step -
accuracy: 0.4667 - loss: 1.2440
Epoch 3: val_loss did not improve from 0.97379
59/59          87s 549ms/step -
accuracy: 0.4665 - loss: 1.2448 - val_accuracy: 0.5536 - val_loss: 1.1431
Epoch 4/50
 1/59          3s 52ms/step - accuracy:
0.5625 - loss: 0.9831
Epoch 4: val_loss did not improve from 0.97379
59/59          0s 487us/step -
accuracy: 0.5625 - loss: 0.9831 - val_accuracy: 0.4286 - val_loss: 1.1856
Epoch 5/50
59/59          0s 474ms/step -
accuracy: 0.5264 - loss: 1.1673
Epoch 5: val_loss did not improve from 0.97379
59/59          36s 545ms/step -
accuracy: 0.5268 - loss: 1.1669 - val_accuracy: 0.5759 - val_loss: 1.1152
Epoch 6/50
 1/59          3s 53ms/step - accuracy:
0.5625 - loss: 1.1106
Epoch 6: val_loss did not improve from 0.97379
59/59          0s 488us/step -
accuracy: 0.5625 - loss: 1.1106 - val_accuracy: 0.5000 - val_loss: 0.9987
Epoch 7/50
59/59          0s 490ms/step -
accuracy: 0.5754 - loss: 1.0838
Epoch 7: val_loss did not improve from 0.97379
59/59          41s 558ms/step -
accuracy: 0.5754 - loss: 1.0836 - val_accuracy: 0.6101 - val_loss: 1.0163

```

```

[ ]: # Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_generator, steps=test_generator.
    ↪samples // batch_size)
print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_accuracy}")

```

```

4/4          1s 217ms/step -
accuracy: 0.3146 - loss: 1.4061
Test Loss: 1.2935124635696411
Test Accuracy: 0.40625

```

```

[ ]: import numpy as np
import matplotlib.pyplot as plt

```

```

import seaborn as sns
from sklearn.metrics import confusion_matrix
from tensorflow.keras.models import load_model

# Generate predictions for the test data without specifying 'steps'
test_predictions = model.predict(test_generator, verbose=1)

# Get the predicted class labels
predicted_labels = np.argmax(test_predictions, axis=1)

# Get the true class labels from the test generator
true_labels = test_generator.classes


# Generate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(10, 7))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=test_generator.
    ↪class_indices.keys(), yticklabels=test_generator.class_indices.keys())
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

```

5/5

2s 319ms/step



As expected it was so hard to know from the first time  
so i will not waste time in tring find the best CNN model  
but i will try **Pretrained models**

### 3.1 Pre-Trained MobileNetV2

```
[ ]: import tensorflow as tf
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout
from tensorflow.keras.models import Model

# Define the number of classes in your dataset
num_classes = 5

# Load the pre-trained MobileNetV2 model
```

```

base_model = MobileNetV2(weights='imagenet', include_top=False,
    ↪input_shape=(224, 224, 3))

# Freeze the base model's layers to prevent them from being trained
base_model.trainable = False

# Add custom classification head
x = base_model.output
x = GlobalAveragePooling2D()(x) # Reduce dimensions to prevent overfitting
x = Dropout(0.3)(x) # Dropout for regularization
x = Dense(128, activation='relu')(x) # Add a fully connected layer
x = Dropout(0.3)(x) # Dropout for regularization
output = Dense(num_classes, activation='softmax')(x) # Final classification
    ↪layer

# Create the model
model = Model(inputs=base_model.input, outputs=output)

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

# Display the model summary
model.summary()

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\\_v2/mobilenet\\_v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_1.0\\_224\\_no\\_top.h5](https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5)

9406464/9406464 2s  
0us/step

Model: "functional\_2"

Layer (type)	Output Shape	Param #	Connected
↪to			
input_layer_2 (InputLayer)	(None, 224, 224, 3)	0	-
↪			
Conv1 (Conv2D)	(None, 112, 112, 32)	864	
↪input_layer_2[0][0]			
bn_Conv1	(None, 112, 112, 32)	128	
↪Conv1[0][0]			



(BatchNormalization)			└
↳			
Conv1_relu (ReLU)	(None, 112, 112, 32)	0	└
↳bn_Conv1[0][0]			
expanded_conv_depthwise	(None, 112, 112, 32)	288	└
↳Conv1_relu[0][0]			
(DepthwiseConv2D)			└
↳			
expanded_conv_depthwise_...	(None, 112, 112, 32)	128	└
↳expanded_conv_depthwi...			
(BatchNormalization)			└
↳			
expanded_conv_depthwise_...	(None, 112, 112, 32)	0	└
↳expanded_conv_depthwi...			
(ReLU)			└
↳			
expanded_conv_project	(None, 112, 112, 16)	512	└
↳expanded_conv_depthwi...			
(Conv2D)			└
↳			
expanded_conv_project_BN	(None, 112, 112, 16)	64	└
↳expanded_conv_project...			
(BatchNormalization)			└
↳			
block_1_expand (Conv2D)	(None, 112, 112, 96)	1,536	└
↳expanded_conv_project...			
block_1_expand_BN	(None, 112, 112, 96)	384	└
↳block_1_expand[0][0]			
(BatchNormalization)			└
↳			
block_1_expand_relu	(None, 112, 112, 96)	0	└
↳block_1_expand_BN[0][...			
(ReLU)			└
↳			
block_1_pad	(None, 113, 113, 96)	0	└
↳block_1_expand_relu[0...			

(ZeroPadding2D)			⌵
↪			
block_1_depthwise	(None, 56, 56, 96)	864	⌵
↪block_1_pad[0][0]			
(DepthwiseConv2D)			⌵
↪			
block_1_depthwise_BN	(None, 56, 56, 96)	384	⌵
↪block_1_depthwise[0][...]			
(BatchNormalization)			⌵
↪			
block_1_depthwise_relu	(None, 56, 56, 96)	0	⌵
↪block_1_depthwise_BN[...]			
(ReLU)			⌵
↪			
block_1_project (Conv2D)	(None, 56, 56, 24)	2,304	⌵
↪block_1_depthwise_relu...			
block_1_project_BN	(None, 56, 56, 24)	96	⌵
↪block_1_project[0][0]			
(BatchNormalization)			⌵
↪			
block_2_expand (Conv2D)	(None, 56, 56, 144)	3,456	⌵
↪block_1_project_BN[0]...			
block_2_expand_BN	(None, 56, 56, 144)	576	⌵
↪block_2_expand[0][0]			
(BatchNormalization)			⌵
↪			
block_2_expand_relu	(None, 56, 56, 144)	0	⌵
↪block_2_expand_BN[0][...]			
(ReLU)			⌵
↪			
block_2_depthwise	(None, 56, 56, 144)	1,296	⌵
↪block_2_expand_relu[0]...			
(DepthwiseConv2D)			⌵
↪			
block_2_depthwise_BN	(None, 56, 56, 144)	576	⌵
↪block_2_depthwise[0][...]			

(BatchNormalization)			└
↪			
block_2_depthwise_relu	(None, 56, 56, 144)	0	└
↪block_2_depthwise_BN[...]			
(ReLU)			└
↪			
block_2_project (Conv2D)	(None, 56, 56, 24)	3,456	└
↪block_2_depthwise_relu...			
block_2_project_BN	(None, 56, 56, 24)	96	└
↪block_2_project[0][0]			
(BatchNormalization)			└
↪			
block_2_add (Add)	(None, 56, 56, 24)	0	└
↪block_1_project_BN[0]...			
			└
↪block_2_project_BN[0]...			
block_3_expand (Conv2D)	(None, 56, 56, 144)	3,456	└
↪block_2_add[0][0]			
block_3_expand_BN	(None, 56, 56, 144)	576	└
↪block_3_expand[0][0]			
(BatchNormalization)			└
↪			
block_3_expand_relu	(None, 56, 56, 144)	0	└
↪block_3_expand_BN[0][...]			
(ReLU)			└
↪			
block_3_pad	(None, 57, 57, 144)	0	└
↪block_3_expand_relu[0]...			
(ZeroPadding2D)			└
↪			
block_3_depthwise	(None, 28, 28, 144)	1,296	└
↪block_3_pad[0][0]			
(DepthwiseConv2D)			└
↪			
block_3_depthwise_BN	(None, 28, 28, 144)	576	└
↪block_3_depthwise[0][...]			

(BatchNormalization)			└
↳			
block_3_depthwise_relu	(None, 28, 28, 144)	0	└
↳block_3_depthwise_BN[...]			
(ReLU)			└
↳			
block_3_project (Conv2D)	(None, 28, 28, 32)	4,608	└
↳block_3_depthwise_relu...			
block_3_project_BN	(None, 28, 28, 32)	128	└
↳block_3_project[0][0]			
(BatchNormalization)			└
↳			
block_4_expand (Conv2D)	(None, 28, 28, 192)	6,144	└
↳block_3_project_BN[0]...			
block_4_expand_BN	(None, 28, 28, 192)	768	└
↳block_4_expand[0][0]			
(BatchNormalization)			└
↳			
block_4_expand_relu	(None, 28, 28, 192)	0	└
↳block_4_expand_BN[0][...]			
(ReLU)			└
↳			
block_4_depthwise	(None, 28, 28, 192)	1,728	└
↳block_4_expand_relu[0...]			
(DepthwiseConv2D)			└
↳			
block_4_depthwise_BN	(None, 28, 28, 192)	768	└
↳block_4_depthwise[0][...]			
(BatchNormalization)			└
↳			
block_4_depthwise_relu	(None, 28, 28, 192)	0	└
↳block_4_depthwise_BN[...]			
(ReLU)			└
↳			
block_4_project (Conv2D)	(None, 28, 28, 32)	6,144	└
↳block_4_depthwise_relu...			

block_4_project_BN ↳block_4_project[0][0] (BatchNormalization)	(None, 28, 28, 32)	128	↳
↳			
block_4_add (Add) ↳block_3_project_BN[0]...	(None, 28, 28, 32)	0	↳
↳block_4_project_BN[0]...			↳
block_5_expand (Conv2D) ↳block_4_add[0][0]	(None, 28, 28, 192)	6,144	↳
block_5_expand_BN ↳block_5_expand[0][0] (BatchNormalization)	(None, 28, 28, 192)	768	↳
↳			
block_5_expand_relu ↳block_5_expand_BN[0][...] (ReLU)	(None, 28, 28, 192)	0	↳
↳			
block_5_depthwise ↳block_5_expand_relu[0...] (DepthwiseConv2D)	(None, 28, 28, 192)	1,728	↳
↳			
block_5_depthwise_BN ↳block_5_depthwise[0][...] (BatchNormalization)	(None, 28, 28, 192)	768	↳
↳			
block_5_depthwise_relu ↳block_5_depthwise_BN[...] (ReLU)	(None, 28, 28, 192)	0	↳
↳			
block_5_project (Conv2D) ↳block_5_depthwise_relu...	(None, 28, 28, 32)	6,144	↳
block_5_project_BN ↳block_5_project[0][0] (BatchNormalization)	(None, 28, 28, 32)	128	↳
↳			

block_5_add (Add)	(None, 28, 28, 32)	0	┐
↳block_4_add[0][0],			
↳block_5_project_BN[0]...			┐
block_6_expand (Conv2D)	(None, 28, 28, 192)	6,144	┐
↳block_5_add[0][0]			
block_6_expand_BN	(None, 28, 28, 192)	768	┐
↳block_6_expand[0][0]			
(BatchNormalization)			┐
↳			
block_6_expand_relu	(None, 28, 28, 192)	0	┐
↳block_6_expand_BN[0][...]			
(ReLU)			┐
↳			
block_6_pad	(None, 29, 29, 192)	0	┐
↳block_6_expand_relu[0...]			
(ZeroPadding2D)			┐
↳			
block_6_depthwise	(None, 14, 14, 192)	1,728	┐
↳block_6_pad[0][0]			
(DepthwiseConv2D)			┐
↳			
block_6_depthwise_BN	(None, 14, 14, 192)	768	┐
↳block_6_depthwise[0][...]			
(BatchNormalization)			┐
↳			
block_6_depthwise_relu	(None, 14, 14, 192)	0	┐
↳block_6_depthwise_BN[...]			
(ReLU)			┐
↳			
block_6_project (Conv2D)	(None, 14, 14, 64)	12,288	┐
↳block_6_depthwise_relu...			
block_6_project_BN	(None, 14, 14, 64)	256	┐
↳block_6_project[0][0]			
(BatchNormalization)			┐
↳			

block_7_expand (Conv2D) ↳block_6_project_BN[0]...	(None, 14, 14, 384)	24,576	⌞
block_7_expand_BN ↳block_7_expand[0][0] (BatchNormalization)	(None, 14, 14, 384)	1,536	⌞
↳			
block_7_expand_relu ↳block_7_expand_BN[0][...] (ReLU)	(None, 14, 14, 384)	0	⌞
↳			
block_7_depthwise ↳block_7_expand_relu[0...] (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	⌞
↳			
block_7_depthwise_BN ↳block_7_depthwise[0][...] (BatchNormalization)	(None, 14, 14, 384)	1,536	⌞
↳			
block_7_depthwise_relu ↳block_7_depthwise_BN[...] (ReLU)	(None, 14, 14, 384)	0	⌞
↳			
block_7_project (Conv2D) ↳block_7_depthwise_rel...	(None, 14, 14, 64)	24,576	⌞
block_7_project_BN ↳block_7_project[0][0] (BatchNormalization)	(None, 14, 14, 64)	256	⌞
↳			
block_7_add (Add) ↳block_6_project_BN[0]...	(None, 14, 14, 64)	0	⌞
↳block_7_project_BN[0]...			⌞
block_8_expand (Conv2D) ↳block_7_add[0][0]	(None, 14, 14, 384)	24,576	⌞

block_8_expand_BN ↳block_8_expand[0][0] (BatchNormalization)	(None, 14, 14, 384)	1,536	┐	┐
↳				
block_8_expand_relu ↳block_8_expand_BN[0][...] (ReLU)	(None, 14, 14, 384)	0	┐	┐
↳				
block_8_depthwise ↳block_8_expand_relu[0...] (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	┐	┐
↳				
block_8_depthwise_BN ↳block_8_depthwise[0][...] (BatchNormalization)	(None, 14, 14, 384)	1,536	┐	┐
↳				
block_8_depthwise_relu ↳block_8_depthwise_BN[...] (ReLU)	(None, 14, 14, 384)	0	┐	┐
↳				
block_8_project (Conv2D) ↳block_8_depthwise_relu...	(None, 14, 14, 64)	24,576	┐	
block_8_project_BN ↳block_8_project[0][0] (BatchNormalization)	(None, 14, 14, 64)	256	┐	┐
↳				
block_8_add (Add) ↳block_7_add[0][0],  ↳block_8_project_BN[0]...	(None, 14, 14, 64)	0	┐	┐
block_9_expand (Conv2D) ↳block_8_add[0][0]	(None, 14, 14, 384)	24,576	┐	
block_9_expand_BN ↳block_9_expand[0][0] (BatchNormalization)	(None, 14, 14, 384)	1,536	┐	┐
↳				



block_9_expand_relu ↳block_9_expand_BN[0] [...] (ReLU)	(None, 14, 14, 384)	0	┐
↳			
block_9_depthwise ↳block_9_expand_relu[0]... (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	┐
↳			
block_9_depthwise_BN ↳block_9_depthwise[0] [...] (BatchNormalization)	(None, 14, 14, 384)	1,536	┐
↳			
block_9_depthwise_relu ↳block_9_depthwise_BN[...] (ReLU)	(None, 14, 14, 384)	0	┐
↳			
block_9_project (Conv2D) ↳block_9_depthwise_rel...	(None, 14, 14, 64)	24,576	┐
block_9_project_BN ↳block_9_project[0][0] (BatchNormalization)	(None, 14, 14, 64)	256	┐
↳			
block_9_add (Add) ↳block_8_add[0][0],  ↳block_9_project_BN[0]...	(None, 14, 14, 64)	0	┐
block_10_expand (Conv2D) ↳block_9_add[0][0]	(None, 14, 14, 384)	24,576	┐
block_10_expand_BN ↳block_10_expand[0][0] (BatchNormalization)	(None, 14, 14, 384)	1,536	┐
↳			
block_10_expand_relu ↳block_10_expand_BN[0]... (ReLU)	(None, 14, 14, 384)	0	┐
↳			

block_10_depthwise ↳block_10_expand_relu[... (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	┐
↳			
block_10_depthwise_BN ↳block_10_depthwise[0]... (BatchNormalization)	(None, 14, 14, 384)	1,536	┐
↳			
block_10_depthwise_relu ↳block_10_depthwise_BN... (ReLU)	(None, 14, 14, 384)	0	┐
↳			
block_10_project (Conv2D) ↳block_10_depthwise_re...	(None, 14, 14, 96)	36,864	┐
block_10_project_BN ↳block_10_project[0][0] (BatchNormalization)	(None, 14, 14, 96)	384	┐
↳			
block_11_expand (Conv2D) ↳block_10_project_BN[0...	(None, 14, 14, 576)	55,296	┐
block_11_expand_BN ↳block_11_expand[0][0] (BatchNormalization)	(None, 14, 14, 576)	2,304	┐
↳			
block_11_expand_relu ↳block_11_expand_BN[0]... (ReLU)	(None, 14, 14, 576)	0	┐
↳			
block_11_depthwise ↳block_11_expand_relu[... (DepthwiseConv2D)	(None, 14, 14, 576)	5,184	┐
↳			
block_11_depthwise_BN ↳block_11_depthwise[0]... (BatchNormalization)	(None, 14, 14, 576)	2,304	┐
↳			

block_11_depthwise_relu ↳block_11_depthwise_BN... (ReLU)	(None, 14, 14, 576)	0	↳
block_11_project (Conv2D) ↳block_11_depthwise_re...	(None, 14, 14, 96)	55,296	↳
block_11_project_BN ↳block_11_project[0][0] (BatchNormalization)	(None, 14, 14, 96)	384	↳
block_11_add (Add) ↳block_10_project_BN[0...	(None, 14, 14, 96)	0	↳
↳block_11_project_BN[0...			↳
block_12_expand (Conv2D) ↳block_11_add[0][0]	(None, 14, 14, 576)	55,296	↳
block_12_expand_BN ↳block_12_expand[0][0] (BatchNormalization)	(None, 14, 14, 576)	2,304	↳
block_12_expand_relu ↳block_12_expand_BN[0]...	(None, 14, 14, 576)	0	↳
(ReLU)			↳
block_12_depthwise ↳block_12_expand_relu[... (DepthwiseConv2D)	(None, 14, 14, 576)	5,184	↳
block_12_depthwise_BN ↳block_12_depthwise[0]...	(None, 14, 14, 576)	2,304	↳
(BatchNormalization)			↳
block_12_depthwise_relu ↳block_12_depthwise_BN... (ReLU)	(None, 14, 14, 576)	0	↳

block_12_project (Conv2D)	(None, 14, 14, 96)	55,296	┐
↳ block_12_depthwise_re...			
block_12_project_BN	(None, 14, 14, 96)	384	┐
↳ block_12_project[0][0]			
(BatchNormalization)			┐
↳			
block_12_add (Add)	(None, 14, 14, 96)	0	┐
↳ block_11_add[0][0],			
			┐
↳ block_12_project_BN[0...			
block_13_expand (Conv2D)	(None, 14, 14, 576)	55,296	┐
↳ block_12_add[0][0]			
block_13_expand_BN	(None, 14, 14, 576)	2,304	┐
↳ block_13_expand[0][0]			
(BatchNormalization)			┐
↳			
block_13_expand_relu	(None, 14, 14, 576)	0	┐
↳ block_13_expand_BN[0]...			
(ReLU)			┐
↳			
block_13_pad	(None, 15, 15, 576)	0	┐
↳ block_13_expand_relu[...			
(ZeroPadding2D)			┐
↳			
block_13_depthwise	(None, 7, 7, 576)	5,184	┐
↳ block_13_pad[0][0]			
(DepthwiseConv2D)			┐
↳			
block_13_depthwise_BN	(None, 7, 7, 576)	2,304	┐
↳ block_13_depthwise[0]...			
(BatchNormalization)			┐
↳			
block_13_depthwise_relu	(None, 7, 7, 576)	0	┐
↳ block_13_depthwise_BN...			
(ReLU)			┐
↳			

block_13_project (Conv2D)	(None, 7, 7, 160)	92,160	┐
↳ block_13_depthwise_re...			
block_13_project_BN	(None, 7, 7, 160)	640	┐
↳ block_13_project[0][0]			
(BatchNormalization)			┐
↳			
block_14_expand (Conv2D)	(None, 7, 7, 960)	153,600	┐
↳ block_13_project_BN[0...			
block_14_expand_BN	(None, 7, 7, 960)	3,840	┐
↳ block_14_expand[0][0]			
(BatchNormalization)			┐
↳			
block_14_expand_relu	(None, 7, 7, 960)	0	┐
↳ block_14_expand_BN[0]...			
(ReLU)			┐
↳			
block_14_depthwise	(None, 7, 7, 960)	8,640	┐
↳ block_14_expand_relu[...			
(DepthwiseConv2D)			┐
↳			
block_14_depthwise_BN	(None, 7, 7, 960)	3,840	┐
↳ block_14_depthwise[0]...			
(BatchNormalization)			┐
↳			
block_14_depthwise_relu	(None, 7, 7, 960)	0	┐
↳ block_14_depthwise_BN...			
(ReLU)			┐
↳			
block_14_project (Conv2D)	(None, 7, 7, 160)	153,600	┐
↳ block_14_depthwise_re...			
block_14_project_BN	(None, 7, 7, 160)	640	┐
↳ block_14_project[0][0]			
(BatchNormalization)			┐
↳			
block_14_add (Add)	(None, 7, 7, 160)	0	┐
↳ block_13_project_BN[0...			

↳block_14_project_BN[0...			↳
block_15_expand (Conv2D)	(None, 7, 7, 960)	153,600	↳
↳block_14_add[0][0]			
block_15_expand_BN	(None, 7, 7, 960)	3,840	↳
↳block_15_expand[0][0]			
(BatchNormalization)			↳
↳			
block_15_expand_relu	(None, 7, 7, 960)	0	↳
↳block_15_expand_BN[0]...			
(ReLU)			↳
↳			
block_15_depthwise	(None, 7, 7, 960)	8,640	↳
↳block_15_expand_relu[...			
(DepthwiseConv2D)			↳
↳			
block_15_depthwise_BN	(None, 7, 7, 960)	3,840	↳
↳block_15_depthwise[0]...			
(BatchNormalization)			↳
↳			
block_15_depthwise_relu	(None, 7, 7, 960)	0	↳
↳block_15_depthwise_BN...			
(ReLU)			↳
↳			
block_15_project (Conv2D)	(None, 7, 7, 160)	153,600	↳
↳block_15_depthwise_re...			
block_15_project_BN	(None, 7, 7, 160)	640	↳
↳block_15_project[0][0]			
(BatchNormalization)			↳
↳			
block_15_add (Add)	(None, 7, 7, 160)	0	↳
↳block_14_add[0][0],			
			↳
↳block_15_project_BN[0...			
block_16_expand (Conv2D)	(None, 7, 7, 960)	153,600	↳
↳block_15_add[0][0]			

block_16_expand_BN ↳ block_16_expand[0][0] (BatchNormalization)	(None, 7, 7, 960)	3,840	↳
block_16_expand_relu ↳ block_16_expand_BN[0]...	(None, 7, 7, 960)	0	↳
block_16_depthwise ↳ block_16_expand_relu[... (DepthwiseConv2D)	(None, 7, 7, 960)	8,640	↳
block_16_depthwise_BN ↳ block_16_depthwise[0]...	(None, 7, 7, 960)	3,840	↳
block_16_depthwise_relu ↳ block_16_depthwise_BN...	(None, 7, 7, 960)	0	↳
block_16_project (Conv2D) ↳ block_16_depthwise_re...	(None, 7, 7, 320)	307,200	↳
block_16_project_BN ↳ block_16_project[0][0] (BatchNormalization)	(None, 7, 7, 320)	1,280	↳
Conv_1 (Conv2D) ↳ block_16_project_BN[0...	(None, 7, 7, 1280)	409,600	↳
Conv_1_bn ↳ Conv_1[0][0] (BatchNormalization)	(None, 7, 7, 1280)	5,120	↳
out_relu (ReLU) ↳ Conv_1_bn[0][0]	(None, 7, 7, 1280)	0	↳

```

global_average_pooling2d... (None, 1280) 0
↳out_relu[0][0]
(GlobalAveragePooling2D)
↳

dropout_4 (Dropout) (None, 1280) 0
↳global_average_poolin...

dense_4 (Dense) (None, 128) 163,968
↳dropout_4[0][0]

dropout_5 (Dropout) (None, 128) 0
↳dense_4[0][0]

dense_5 (Dense) (None, 5) 645
↳dropout_5[0][0]

```

Total params: 2,422,597 (9.24 MB)

Trainable params: 164,613 (643.02 KB)

Non-trainable params: 2,257,984 (8.61 MB)

```

[ ]: from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_dir = '/content/drive/MyDrive/Flowers_Divided/train' # Path to your
↳train data
val_dir = '/content/drive/MyDrive/Flowers_Divided/val' # Path to your
↳validation data

# Create ImageDataGenerators
train_datagen = ImageDataGenerator(rescale=1./255, rotation_range=20,
↳width_shift_range=0.2, height_shift_range=0.2, horizontal_flip=True)
val_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(train_dir,
↳target_size=(224, 224), batch_size=32, class_mode='categorical')
val_generator = val_datagen.flow_from_directory(val_dir, target_size=(224,
↳224), batch_size=32, class_mode='categorical')

```

Found 1919 images belonging to 5 classes.

Found 686 images belonging to 5 classes.



```
[ ]: # Train the model
history = model.fit(train_generator, epochs=10, validation_data=val_generator)
```

Epoch 1/10

/usr/local/lib/python3.10/dist-

packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:122:

UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self.\_warn\_if\_super\_not\_called()

60/60 57s 730ms/step -

accuracy: 0.5499 - loss: 1.1626 - val\_accuracy: 0.8397 - val\_loss: 0.4573

Epoch 2/10

60/60 35s 523ms/step -

accuracy: 0.7890 - loss: 0.5565 - val\_accuracy: 0.8703 - val\_loss: 0.3778

Epoch 3/10

60/60 35s 531ms/step -

accuracy: 0.8240 - loss: 0.4652 - val\_accuracy: 0.8557 - val\_loss: 0.3841

Epoch 4/10

60/60 34s 512ms/step -

accuracy: 0.8439 - loss: 0.3984 - val\_accuracy: 0.8907 - val\_loss: 0.3281

Epoch 5/10

60/60 41s 521ms/step -

accuracy: 0.8489 - loss: 0.4177 - val\_accuracy: 0.8717 - val\_loss: 0.3581

Epoch 6/10

60/60 36s 536ms/step -

accuracy: 0.8642 - loss: 0.3809 - val\_accuracy: 0.8950 - val\_loss: 0.3194

Epoch 7/10

60/60 40s 528ms/step -

accuracy: 0.8739 - loss: 0.3457 - val\_accuracy: 0.8950 - val\_loss: 0.3073

Epoch 8/10

60/60 34s 515ms/step -

accuracy: 0.8950 - loss: 0.2812 - val\_accuracy: 0.8965 - val\_loss: 0.3214

Epoch 9/10

60/60 36s 543ms/step -

accuracy: 0.9021 - loss: 0.2899 - val\_accuracy: 0.8834 - val\_loss: 0.3295

Epoch 10/10

60/60 34s 520ms/step -

accuracy: 0.9072 - loss: 0.2762 - val\_accuracy: 0.8980 - val\_loss: 0.2977

```
[ ]: # Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_generator, steps=test_generator.
    ↪samples // batch_size)
print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_accuracy}")
```

4/4                    1s 254ms/step -  
accuracy: 0.8615 - loss: 0.4090  
Test Loss: 0.44765931367874146  
Test Accuracy: 0.8515625

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
from tensorflow.keras.models import load_model

# Generate predictions for the test data without specifying 'steps'
test_predictions = model.predict(test_generator, verbose=1)


# Get the predicted class labels
predicted_labels = np.argmax(test_predictions, axis=1)

# Get the true class labels from the test generator
true_labels = test_generator.classes

# Generate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(10, 7))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=test_generator.
    class_indices.keys(), yticklabels=test_generator.class_indices.keys())
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

5/5                    8s 1s/step



Nice one

```
[ ]: model.save("/content/drive/MyDrive/Flowers_Divided/MobileNetV2.keras")
```

### 3.2 fine-tune the MobileNetV2 model

**Steps for Fine-tuning MobileNetV2** Train the Custom Head: Initially, freeze the base model and train only the custom head (already done in the previous step).

Unfreeze Layers: Unfreeze some of the deeper layers of the base model while keeping the earlier layers frozen (to preserve general features).

Lower Learning Rate: Use a smaller learning rate for fine-tuning to avoid destroying the pre-trained weights.

Re-train the Model: Train the model with the updated trainable layers.

```
[ ]: # Unfreeze some layers of the base model
base_model.trainable = True # Unfreeze the entire model for fine-tuning
    ↳ (optional: unfreeze specific layers)
```

```

# Optionally, unfreeze specific layers (e.g., last 20 layers)
for layer in base_model.layers[:-20]: # Freeze all layers except the last 20
    layer.trainable = False

# Compile the model with a lower learning rate
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4), # Lower
              ↪learning rate
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()

```

Model: "functional\_1"

Layer (type) ↪to	Output Shape	Param #	Connected
input_layer_1 (InputLayer) ↪	(None, 224, 224, 3)	0	-
Conv1 (Conv2D) ↪input_layer_1[0][0]	(None, 112, 112, 32)	864	
bn_Conv1 ↪Conv1[0][0] (BatchNormalization) ↪	(None, 112, 112, 32)	128	
Conv1_relu (ReLU) ↪bn_Conv1[0][0]	(None, 112, 112, 32)	0	
expanded_conv_depthwise ↪Conv1_relu[0][0] (DepthwiseConv2D) ↪	(None, 112, 112, 32)	288	
expanded_conv_depthwise_... ↪expanded_conv_depthwi... (BatchNormalization) ↪	(None, 112, 112, 32)	128	
expanded_conv_depthwise_... ↪expanded_conv_depthwi...	(None, 112, 112, 32)	0	

(ReLU)			└
↪			
expanded_conv_project	(None, 112, 112, 16)	512	└
↪expanded_conv_depthwi...			
(Conv2D)			└
↪			
expanded_conv_project_BN	(None, 112, 112, 16)	64	└
↪expanded_conv_project...			
(BatchNormalization)			└
↪			
block_1_expand (Conv2D)	(None, 112, 112, 96)	1,536	└
↪expanded_conv_project...			
block_1_expand_BN	(None, 112, 112, 96)	384	└
↪block_1_expand[0][0]			
(BatchNormalization)			└
↪			
block_1_expand_relu	(None, 112, 112, 96)	0	└
↪block_1_expand_BN[0][...			
(ReLU)			└
↪			
block_1_pad	(None, 113, 113, 96)	0	└
↪block_1_expand_relu[0...			
(ZeroPadding2D)			└
↪			
block_1_depthwise	(None, 56, 56, 96)	864	└
↪block_1_pad[0][0]			
(DepthwiseConv2D)			└
↪			
block_1_depthwise_BN	(None, 56, 56, 96)	384	└
↪block_1_depthwise[0][...			
(BatchNormalization)			└
↪			
block_1_depthwise_relu	(None, 56, 56, 96)	0	└
↪block_1_depthwise_BN[...			
(ReLU)			└
↪			

block_1_project (Conv2D)	(None, 56, 56, 24)	2,304	⌞
↳ block_1_depthwise_rel...			
block_1_project_BN	(None, 56, 56, 24)	96	⌞
↳ block_1_project[0][0]			
(BatchNormalization)			⌞
↳			
block_2_expand (Conv2D)	(None, 56, 56, 144)	3,456	⌞
↳ block_1_project_BN[0]...			
block_2_expand_BN	(None, 56, 56, 144)	576	⌞
↳ block_2_expand[0][0]			
(BatchNormalization)			⌞
↳			
block_2_expand_relu	(None, 56, 56, 144)	0	⌞
↳ block_2_expand_BN[0][...]			
(ReLU)			⌞
↳			
block_2_depthwise	(None, 56, 56, 144)	1,296	⌞
↳ block_2_expand_relu[0]...			
(DepthwiseConv2D)			⌞
↳			
block_2_depthwise_BN	(None, 56, 56, 144)	576	⌞
↳ block_2_depthwise[0][...]			
(BatchNormalization)			⌞
↳			
block_2_depthwise_relu	(None, 56, 56, 144)	0	⌞
↳ block_2_depthwise_BN[...]			
(ReLU)			⌞
↳			
block_2_project (Conv2D)	(None, 56, 56, 24)	3,456	⌞
↳ block_2_depthwise_rel...			
block_2_project_BN	(None, 56, 56, 24)	96	⌞
↳ block_2_project[0][0]			
(BatchNormalization)			⌞
↳			
block_2_add (Add)	(None, 56, 56, 24)	0	⌞
↳ block_1_project_BN[0]...			

↳block_2_project_BN[0]...			↳
block_3_expand (Conv2D)	(None, 56, 56, 144)	3,456	↳
↳block_2_add[0][0]			
block_3_expand_BN	(None, 56, 56, 144)	576	↳
↳block_3_expand[0][0]			
(BatchNormalization)			↳
↳			
block_3_expand_relu	(None, 56, 56, 144)	0	↳
↳block_3_expand_BN[0][...]			
(ReLU)			↳
↳			
block_3_pad	(None, 57, 57, 144)	0	↳
↳block_3_expand_relu[0...]			
(ZeroPadding2D)			↳
↳			
block_3_depthwise	(None, 28, 28, 144)	1,296	↳
↳block_3_pad[0][0]			
(DepthwiseConv2D)			↳
↳			
block_3_depthwise_BN	(None, 28, 28, 144)	576	↳
↳block_3_depthwise[0][...]			
(BatchNormalization)			↳
↳			
block_3_depthwise_relu	(None, 28, 28, 144)	0	↳
↳block_3_depthwise_BN[...]			
(ReLU)			↳
↳			
block_3_project (Conv2D)	(None, 28, 28, 32)	4,608	↳
↳block_3_depthwise_relu...			
block_3_project_BN	(None, 28, 28, 32)	128	↳
↳block_3_project[0][0]			
(BatchNormalization)			↳
↳			
block_4_expand (Conv2D)	(None, 28, 28, 192)	6,144	↳
↳block_3_project_BN[0]...			

block_4_expand_BN ↳block_4_expand[0][0] (BatchNormalization)	(None, 28, 28, 192)	768	↳
↳			
block_4_expand_relu ↳block_4_expand_BN[0][...] (ReLU)	(None, 28, 28, 192)	0	↳
↳			
block_4_depthwise ↳block_4_expand_relu[0...] (DepthwiseConv2D)	(None, 28, 28, 192)	1,728	↳
↳			
block_4_depthwise_BN ↳block_4_depthwise[0][...] (BatchNormalization)	(None, 28, 28, 192)	768	↳
↳			
block_4_depthwise_relu ↳block_4_depthwise_BN[...] (ReLU)	(None, 28, 28, 192)	0	↳
↳			
block_4_project (Conv2D) ↳block_4_depthwise_relu...	(None, 28, 28, 32)	6,144	↳
block_4_project_BN ↳block_4_project[0][0] (BatchNormalization)	(None, 28, 28, 32)	128	↳
↳			
block_4_add (Add) ↳block_3_project_BN[0]...	(None, 28, 28, 32)	0	↳
↳block_4_project_BN[0]...			↳
block_5_expand (Conv2D) ↳block_4_add[0][0]	(None, 28, 28, 192)	6,144	↳
block_5_expand_BN ↳block_5_expand[0][0] (BatchNormalization)	(None, 28, 28, 192)	768	↳
↳			



block_5_expand_relu ↳block_5_expand_BN[0][...] (ReLU)	(None, 28, 28, 192)	0	↳
block_5_depthwise ↳block_5_expand_relu[0...] (DepthwiseConv2D)	(None, 28, 28, 192)	1,728	↳
block_5_depthwise_BN ↳block_5_depthwise[0][...] (BatchNormalization)	(None, 28, 28, 192)	768	↳
block_5_depthwise_relu ↳block_5_depthwise_BN[...] (ReLU)	(None, 28, 28, 192)	0	↳
block_5_project (Conv2D) ↳block_5_depthwise_relu...	(None, 28, 28, 32)	6,144	↳
block_5_project_BN ↳block_5_project[0][0] (BatchNormalization)	(None, 28, 28, 32)	128	↳
block_5_add (Add) ↳block_4_add[0][0],  ↳block_5_project_BN[0]...	(None, 28, 28, 32)	0	↳
block_6_expand (Conv2D) ↳block_5_add[0][0]	(None, 28, 28, 192)	6,144	↳
block_6_expand_BN ↳block_6_expand[0][0] (BatchNormalization)	(None, 28, 28, 192)	768	↳
block_6_expand_relu ↳block_6_expand_BN[0][...] (ReLU)	(None, 28, 28, 192)	0	↳

block_6_pad ↳ block_6_expand_relu[0... (ZeroPadding2D)	(None, 29, 29, 192)	0	↳
block_6_depthwise ↳ block_6_pad[0][0] (DepthwiseConv2D)	(None, 14, 14, 192)	1,728	↳
block_6_depthwise_BN ↳ block_6_depthwise[0][... (BatchNormalization)	(None, 14, 14, 192)	768	↳
block_6_depthwise_relu ↳ block_6_depthwise_BN[... (ReLU)	(None, 14, 14, 192)	0	↳
block_6_project (Conv2D) ↳ block_6_depthwise_relu...	(None, 14, 14, 64)	12,288	↳
block_6_project_BN ↳ block_6_project[0][0] (BatchNormalization)	(None, 14, 14, 64)	256	↳
block_7_expand (Conv2D) ↳ block_6_project_BN[0]...	(None, 14, 14, 384)	24,576	↳
block_7_expand_BN ↳ block_7_expand[0][0] (BatchNormalization)	(None, 14, 14, 384)	1,536	↳
block_7_expand_relu ↳ block_7_expand_BN[0][... (ReLU)	(None, 14, 14, 384)	0	↳
block_7_depthwise ↳ block_7_expand_relu[0... (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	↳

block_7_depthwise_BN ↳block_7_depthwise[0][... (BatchNormalization)	(None, 14, 14, 384)	1,536	↳
block_7_depthwise_relu ↳block_7_depthwise_BN[... (ReLU)	(None, 14, 14, 384)	0	↳
block_7_project (Conv2D) ↳block_7_depthwise_relu...	(None, 14, 14, 64)	24,576	↳
block_7_project_BN ↳block_7_project[0][0] (BatchNormalization)	(None, 14, 14, 64)	256	↳
block_7_add (Add) ↳block_6_project_BN[0]...	(None, 14, 14, 64)	0	↳
↳block_7_project_BN[0]...			↳
block_8_expand (Conv2D) ↳block_7_add[0][0]	(None, 14, 14, 384)	24,576	↳
block_8_expand_BN ↳block_8_expand[0][0] (BatchNormalization)	(None, 14, 14, 384)	1,536	↳
block_8_expand_relu ↳block_8_expand_BN[0][... (ReLU)	(None, 14, 14, 384)	0	↳
block_8_depthwise ↳block_8_expand_relu[0... (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	↳
block_8_depthwise_BN ↳block_8_depthwise[0][... (BatchNormalization)	(None, 14, 14, 384)	1,536	↳

block_8_depthwise_relu ↳ block_8_depthwise_BN[...] (ReLU)	(None, 14, 14, 384)	0	↳
↳			
block_8_project (Conv2D) ↳ block_8_depthwise_relu...	(None, 14, 14, 64)	24,576	↳
block_8_project_BN ↳ block_8_project[0][0] (BatchNormalization)	(None, 14, 14, 64)	256	↳
↳			
block_8_add (Add) ↳ block_7_add[0][0],  ↳ block_8_project_BN[0]...	(None, 14, 14, 64)	0	↳
			↳
block_9_expand (Conv2D) ↳ block_8_add[0][0]	(None, 14, 14, 384)	24,576	↳
block_9_expand_BN ↳ block_9_expand[0][0] (BatchNormalization)	(None, 14, 14, 384)	1,536	↳
↳			
block_9_expand_relu ↳ block_9_expand_BN[0][...] (ReLU)	(None, 14, 14, 384)	0	↳
↳			
block_9_depthwise ↳ block_9_expand_relu[0...] (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	↳
↳			
block_9_depthwise_BN ↳ block_9_depthwise[0][...] (BatchNormalization)	(None, 14, 14, 384)	1,536	↳
↳			
block_9_depthwise_relu ↳ block_9_depthwise_BN[...] (ReLU)	(None, 14, 14, 384)	0	↳
↳			

block_9_project (Conv2D) ↳block_9_depthwise_rel...	(None, 14, 14, 64)	24,576	⌞
block_9_project_BN ↳block_9_project[0][0] (BatchNormalization) ↳	(None, 14, 14, 64)	256	⌞
block_9_add (Add) ↳block_8_add[0][0],  ↳block_9_project_BN[0]...	(None, 14, 14, 64)	0	⌞
block_10_expand (Conv2D) ↳block_9_add[0][0]	(None, 14, 14, 384)	24,576	⌞
block_10_expand_BN ↳block_10_expand[0][0] (BatchNormalization) ↳	(None, 14, 14, 384)	1,536	⌞
block_10_expand_relu ↳block_10_expand_BN[0]... (ReLU) ↳	(None, 14, 14, 384)	0	⌞
block_10_depthwise ↳block_10_expand_relu[... (DepthwiseConv2D) ↳	(None, 14, 14, 384)	3,456	⌞
block_10_depthwise_BN ↳block_10_depthwise[0]... (BatchNormalization) ↳	(None, 14, 14, 384)	1,536	⌞
block_10_depthwise_relu ↳block_10_depthwise_BN... (ReLU) ↳	(None, 14, 14, 384)	0	⌞
block_10_project (Conv2D) ↳block_10_depthwise_re...	(None, 14, 14, 96)	36,864	⌞

block_10_project_BN ↳block_10_project[0][0] (BatchNormalization)	(None, 14, 14, 96)	384	↳
block_11_expand (Conv2D) ↳block_10_project_BN[0]...	(None, 14, 14, 576)	55,296	↳
block_11_expand_BN ↳block_11_expand[0][0] (BatchNormalization)	(None, 14, 14, 576)	2,304	↳
block_11_expand_relu ↳block_11_expand_BN[0]...	(None, 14, 14, 576)	0	↳
block_11_depthwise ↳block_11_expand_relu[...] (DepthwiseConv2D)	(None, 14, 14, 576)	5,184	↳
block_11_depthwise_BN ↳block_11_depthwise[0]...	(None, 14, 14, 576)	2,304	↳
block_11_depthwise_relu ↳block_11_depthwise_BN...	(None, 14, 14, 576)	0	↳
block_11_project (Conv2D) ↳block_11_depthwise_re...	(None, 14, 14, 96)	55,296	↳
block_11_project_BN ↳block_11_project[0][0] (BatchNormalization)	(None, 14, 14, 96)	384	↳
block_11_add (Add) ↳block_10_project_BN[0]...	(None, 14, 14, 96)	0	↳
↳block_11_project_BN[0]...			↳

block_12_expand (Conv2D) ↳block_11_add[0][0]	(None, 14, 14, 576)	55,296	┐
block_12_expand_BN ↳block_12_expand[0][0] (BatchNormalization)	(None, 14, 14, 576)	2,304	┐
block_12_expand_relu ↳block_12_expand_BN[0]...	(None, 14, 14, 576)	0	┐
block_12_depthwise ↳block_12_expand_relu[...] (DepthwiseConv2D)	(None, 14, 14, 576)	5,184	┐
block_12_depthwise_BN ↳block_12_depthwise[0]...	(None, 14, 14, 576)	2,304	┐
block_12_depthwise_relu ↳block_12_depthwise_BN...	(None, 14, 14, 576)	0	┐
block_12_project (Conv2D) ↳block_12_depthwise_re...	(None, 14, 14, 96)	55,296	┐
block_12_project_BN ↳block_12_project[0][0] (BatchNormalization)	(None, 14, 14, 96)	384	┐
block_12_add (Add) ↳block_11_add[0][0],  ↳block_12_project_BN[0...	(None, 14, 14, 96)	0	┐
block_13_expand (Conv2D) ↳block_12_add[0][0]	(None, 14, 14, 576)	55,296	┐
block_13_expand_BN ↳block_13_expand[0][0]	(None, 14, 14, 576)	2,304	┐

(BatchNormalization)			└
↪			
block_13_expand_relu	(None, 14, 14, 576)	0	└
↪block_13_expand_BN[0]...			
(ReLU)			└
↪			
block_13_pad	(None, 15, 15, 576)	0	└
↪block_13_expand_relu[...			
(ZeroPadding2D)			└
↪			
block_13_depthwise	(None, 7, 7, 576)	5,184	└
↪block_13_pad[0][0]			
(DepthwiseConv2D)			└
↪			
block_13_depthwise_BN	(None, 7, 7, 576)	2,304	└
↪block_13_depthwise[0]...			
(BatchNormalization)			└
↪			
block_13_depthwise_relu	(None, 7, 7, 576)	0	└
↪block_13_depthwise_BN...			
(ReLU)			└
↪			
block_13_project (Conv2D)	(None, 7, 7, 160)	92,160	└
↪block_13_depthwise_re...			
block_13_project_BN	(None, 7, 7, 160)	640	└
↪block_13_project[0][0]			
(BatchNormalization)			└
↪			
block_14_expand (Conv2D)	(None, 7, 7, 960)	153,600	└
↪block_13_project_BN[0...			
block_14_expand_BN	(None, 7, 7, 960)	3,840	└
↪block_14_expand[0][0]			
(BatchNormalization)			└
↪			
block_14_expand_relu	(None, 7, 7, 960)	0	└
↪block_14_expand_BN[0]...			



(ReLU)			└
↳			
block_14_depthwise	(None, 7, 7, 960)	8,640	└
↳block_14_expand_relu[...]			
(DepthwiseConv2D)			└
↳			
block_14_depthwise_BN	(None, 7, 7, 960)	3,840	└
↳block_14_depthwise[0]...			
(BatchNormalization)			└
↳			
block_14_depthwise_relu	(None, 7, 7, 960)	0	└
↳block_14_depthwise_BN...			
(ReLU)			└
↳			
block_14_project (Conv2D)	(None, 7, 7, 160)	153,600	└
↳block_14_depthwise_re...			
block_14_project_BN	(None, 7, 7, 160)	640	└
↳block_14_project[0][0]			
(BatchNormalization)			└
↳			
block_14_add (Add)	(None, 7, 7, 160)	0	└
↳block_13_project_BN[0...			
			└
↳block_14_project_BN[0...			
block_15_expand (Conv2D)	(None, 7, 7, 960)	153,600	└
↳block_14_add[0][0]			
block_15_expand_BN	(None, 7, 7, 960)	3,840	└
↳block_15_expand[0][0]			
(BatchNormalization)			└
↳			
block_15_expand_relu	(None, 7, 7, 960)	0	└
↳block_15_expand_BN[0]...			
(ReLU)			└
↳			
block_15_depthwise	(None, 7, 7, 960)	8,640	└
↳block_15_expand_relu[...			

(DepthwiseConv2D)			⌞
↪			
block_15_depthwise_BN	(None, 7, 7, 960)	3,840	⌞
↪block_15_depthwise[0]...			
(BatchNormalization)			⌞
↪			
block_15_depthwise_relu	(None, 7, 7, 960)	0	⌞
↪block_15_depthwise_BN...			
(ReLU)			⌞
↪			
block_15_project (Conv2D)	(None, 7, 7, 160)	153,600	⌞
↪block_15_depthwise_re...			
block_15_project_BN	(None, 7, 7, 160)	640	⌞
↪block_15_project[0][0]			
(BatchNormalization)			⌞
↪			
block_15_add (Add)	(None, 7, 7, 160)	0	⌞
↪block_14_add[0][0],			
			⌞
↪block_15_project_BN[0...			
block_16_expand (Conv2D)	(None, 7, 7, 960)	153,600	⌞
↪block_15_add[0][0]			
block_16_expand_BN	(None, 7, 7, 960)	3,840	⌞
↪block_16_expand[0][0]			
(BatchNormalization)			⌞
↪			
block_16_expand_relu	(None, 7, 7, 960)	0	⌞
↪block_16_expand_BN[0]...			
(ReLU)			⌞
↪			
block_16_depthwise	(None, 7, 7, 960)	8,640	⌞
↪block_16_expand_relu[...			
(DepthwiseConv2D)			⌞
↪			
block_16_depthwise_BN	(None, 7, 7, 960)	3,840	⌞
↪block_16_depthwise[0]...			

(BatchNormalization)			└
↳			
block_16_depthwise_relu	(None, 7, 7, 960)	0	└
↳block_16_depthwise_BN...			
(ReLU)			└
↳			
block_16_project (Conv2D)	(None, 7, 7, 320)	307,200	└
↳block_16_depthwise_re...			
block_16_project_BN	(None, 7, 7, 320)	1,280	└
↳block_16_project[0][0]			
(BatchNormalization)			└
↳			
Conv_1 (Conv2D)	(None, 7, 7, 1280)	409,600	└
↳block_16_project_BN[0...			
Conv_1_bn	(None, 7, 7, 1280)	5,120	└
↳Conv_1[0][0]			
(BatchNormalization)			└
↳			
out_relu (ReLU)	(None, 7, 7, 1280)	0	└
↳Conv_1_bn[0][0]			
global_average_pooling2d	(None, 1280)	0	└
↳out_relu[0][0]			
(GlobalAveragePooling2D)			└
↳			
dropout (Dropout)	(None, 1280)	0	└
↳global_average_poolin...			
dense_2 (Dense)	(None, 128)	163,968	└
↳dropout[0][0]			
dropout_1 (Dropout)	(None, 128)	0	└
↳dense_2[0][0]			
dense_3 (Dense)	(None, 5)	645	└
↳dropout_1[0][0]			

Total params: 2,422,597 (9.24 MB)

Trainable params: 1,370,693 (5.23 MB)

Non-trainable params: 1,051,904 (4.01 MB)

```
[ ]: for i, layer in enumerate(model.layers):  
      print(f"Layer {i}: {layer.name}, Trainable: {layer.trainable}")
```

```
Layer 0: input_layer_1, Trainable: False  
Layer 1: Conv1, Trainable: False  
Layer 2: bn_Conv1, Trainable: False  
Layer 3: Conv1_relu, Trainable: False  
Layer 4: expanded_conv_depthwise, Trainable: False  
Layer 5: expanded_conv_depthwise_BN, Trainable: False  
Layer 6: expanded_conv_depthwise_relu, Trainable: False  
Layer 7: expanded_conv_project, Trainable: False  
Layer 8: expanded_conv_project_BN, Trainable: False  
Layer 9: block_1_expand, Trainable: False  
Layer 10: block_1_expand_BN, Trainable: False  
Layer 11: block_1_expand_relu, Trainable: False  
Layer 12: block_1_pad, Trainable: False  
Layer 13: block_1_depthwise, Trainable: False  
Layer 14: block_1_depthwise_BN, Trainable: False  
Layer 15: block_1_depthwise_relu, Trainable: False  
Layer 16: block_1_project, Trainable: False  
Layer 17: block_1_project_BN, Trainable: False  
Layer 18: block_2_expand, Trainable: False  
Layer 19: block_2_expand_BN, Trainable: False  
Layer 20: block_2_expand_relu, Trainable: False  
Layer 21: block_2_depthwise, Trainable: False  
Layer 22: block_2_depthwise_BN, Trainable: False  
Layer 23: block_2_depthwise_relu, Trainable: False  
Layer 24: block_2_project, Trainable: False  
Layer 25: block_2_project_BN, Trainable: False  
Layer 26: block_2_add, Trainable: False  
Layer 27: block_3_expand, Trainable: False  
Layer 28: block_3_expand_BN, Trainable: False  
Layer 29: block_3_expand_relu, Trainable: False  
Layer 30: block_3_pad, Trainable: False  
Layer 31: block_3_depthwise, Trainable: False  
Layer 32: block_3_depthwise_BN, Trainable: False  
Layer 33: block_3_depthwise_relu, Trainable: False  
Layer 34: block_3_project, Trainable: False  
Layer 35: block_3_project_BN, Trainable: False  
Layer 36: block_4_expand, Trainable: False  
Layer 37: block_4_expand_BN, Trainable: False
```

Layer 38: block\_4\_expand\_relu, Trainable: False  
Layer 39: block\_4\_depthwise, Trainable: False  
Layer 40: block\_4\_depthwise\_BN, Trainable: False  
Layer 41: block\_4\_depthwise\_relu, Trainable: False  
Layer 42: block\_4\_project, Trainable: False  
Layer 43: block\_4\_project\_BN, Trainable: False  
Layer 44: block\_4\_add, Trainable: False  
Layer 45: block\_5\_expand, Trainable: False  
Layer 46: block\_5\_expand\_BN, Trainable: False  
Layer 47: block\_5\_expand\_relu, Trainable: False  
Layer 48: block\_5\_depthwise, Trainable: False  
Layer 49: block\_5\_depthwise\_BN, Trainable: False  
Layer 50: block\_5\_depthwise\_relu, Trainable: False  
Layer 51: block\_5\_project, Trainable: False  
Layer 52: block\_5\_project\_BN, Trainable: False  
Layer 53: block\_5\_add, Trainable: False  
Layer 54: block\_6\_expand, Trainable: False  
Layer 55: block\_6\_expand\_BN, Trainable: False  
Layer 56: block\_6\_expand\_relu, Trainable: False  
Layer 57: block\_6\_pad, Trainable: False  
Layer 58: block\_6\_depthwise, Trainable: False  
Layer 59: block\_6\_depthwise\_BN, Trainable: False  
Layer 60: block\_6\_depthwise\_relu, Trainable: False  
Layer 61: block\_6\_project, Trainable: False  
Layer 62: block\_6\_project\_BN, Trainable: False  
Layer 63: block\_7\_expand, Trainable: False  
Layer 64: block\_7\_expand\_BN, Trainable: False  
Layer 65: block\_7\_expand\_relu, Trainable: False  
Layer 66: block\_7\_depthwise, Trainable: False  
Layer 67: block\_7\_depthwise\_BN, Trainable: False  
Layer 68: block\_7\_depthwise\_relu, Trainable: False  
Layer 69: block\_7\_project, Trainable: False  
Layer 70: block\_7\_project\_BN, Trainable: False  
Layer 71: block\_7\_add, Trainable: False  
Layer 72: block\_8\_expand, Trainable: False  
Layer 73: block\_8\_expand\_BN, Trainable: False  
Layer 74: block\_8\_expand\_relu, Trainable: False  
Layer 75: block\_8\_depthwise, Trainable: False  
Layer 76: block\_8\_depthwise\_BN, Trainable: False  
Layer 77: block\_8\_depthwise\_relu, Trainable: False  
Layer 78: block\_8\_project, Trainable: False  
Layer 79: block\_8\_project\_BN, Trainable: False  
Layer 80: block\_8\_add, Trainable: False  
Layer 81: block\_9\_expand, Trainable: False  
Layer 82: block\_9\_expand\_BN, Trainable: False  
Layer 83: block\_9\_expand\_relu, Trainable: False  
Layer 84: block\_9\_depthwise, Trainable: False  
Layer 85: block\_9\_depthwise\_BN, Trainable: False

Layer 86: block\_9\_depthwise\_relu, Trainable: False  
Layer 87: block\_9\_project, Trainable: False  
Layer 88: block\_9\_project\_BN, Trainable: False  
Layer 89: block\_9\_add, Trainable: False  
Layer 90: block\_10\_expand, Trainable: False  
Layer 91: block\_10\_expand\_BN, Trainable: False  
Layer 92: block\_10\_expand\_relu, Trainable: False  
Layer 93: block\_10\_depthwise, Trainable: False  
Layer 94: block\_10\_depthwise\_BN, Trainable: False  
Layer 95: block\_10\_depthwise\_relu, Trainable: False  
Layer 96: block\_10\_project, Trainable: False  
Layer 97: block\_10\_project\_BN, Trainable: False  
Layer 98: block\_11\_expand, Trainable: False  
Layer 99: block\_11\_expand\_BN, Trainable: False  
Layer 100: block\_11\_expand\_relu, Trainable: False  
Layer 101: block\_11\_depthwise, Trainable: False  
Layer 102: block\_11\_depthwise\_BN, Trainable: False  
Layer 103: block\_11\_depthwise\_relu, Trainable: False  
Layer 104: block\_11\_project, Trainable: False  
Layer 105: block\_11\_project\_BN, Trainable: False  
Layer 106: block\_11\_add, Trainable: False  
Layer 107: block\_12\_expand, Trainable: False  
Layer 108: block\_12\_expand\_BN, Trainable: False  
Layer 109: block\_12\_expand\_relu, Trainable: False  
Layer 110: block\_12\_depthwise, Trainable: False  
Layer 111: block\_12\_depthwise\_BN, Trainable: False  
Layer 112: block\_12\_depthwise\_relu, Trainable: False  
Layer 113: block\_12\_project, Trainable: False  
Layer 114: block\_12\_project\_BN, Trainable: False  
Layer 115: block\_12\_add, Trainable: False  
Layer 116: block\_13\_expand, Trainable: False  
Layer 117: block\_13\_expand\_BN, Trainable: False  
Layer 118: block\_13\_expand\_relu, Trainable: False  
Layer 119: block\_13\_pad, Trainable: False  
Layer 120: block\_13\_depthwise, Trainable: False  
Layer 121: block\_13\_depthwise\_BN, Trainable: False  
Layer 122: block\_13\_depthwise\_relu, Trainable: False  
Layer 123: block\_13\_project, Trainable: False  
Layer 124: block\_13\_project\_BN, Trainable: False  
Layer 125: block\_14\_expand, Trainable: False  
Layer 126: block\_14\_expand\_BN, Trainable: False  
Layer 127: block\_14\_expand\_relu, Trainable: False  
Layer 128: block\_14\_depthwise, Trainable: False  
Layer 129: block\_14\_depthwise\_BN, Trainable: False  
Layer 130: block\_14\_depthwise\_relu, Trainable: False  
Layer 131: block\_14\_project, Trainable: False  
Layer 132: block\_14\_project\_BN, Trainable: False  
Layer 133: block\_14\_add, Trainable: False

```

Layer 134: block_15_expand, Trainable: True
Layer 135: block_15_expand_BN, Trainable: True
Layer 136: block_15_expand_relu, Trainable: True
Layer 137: block_15_depthwise, Trainable: True
Layer 138: block_15_depthwise_BN, Trainable: True
Layer 139: block_15_depthwise_relu, Trainable: True
Layer 140: block_15_project, Trainable: True
Layer 141: block_15_project_BN, Trainable: True
Layer 142: block_15_add, Trainable: True
Layer 143: block_16_expand, Trainable: True
Layer 144: block_16_expand_BN, Trainable: True
Layer 145: block_16_expand_relu, Trainable: True
Layer 146: block_16_depthwise, Trainable: True
Layer 147: block_16_depthwise_BN, Trainable: True
Layer 148: block_16_depthwise_relu, Trainable: True
Layer 149: block_16_project, Trainable: True
Layer 150: block_16_project_BN, Trainable: True
Layer 151: Conv_1, Trainable: True
Layer 152: Conv_1_bn, Trainable: True
Layer 153: out_relu, Trainable: True
Layer 154: global_average_pooling2d, Trainable: True
Layer 155: dropout, Trainable: True
Layer 156: dense_2, Trainable: True
Layer 157: dropout_1, Trainable: True
Layer 158: dense_3, Trainable: True

```

```

[ ]: # Re-train the model
      history_fine_tune = model.fit(train_generator,
                                   epochs=10, # Add more epochs for fine-tuning
                                   validation_data=val_generator)

```

```

Epoch 1/10
60/60          65s 791ms/step -
accuracy: 0.8386 - loss: 0.4489 - val_accuracy: 0.8586 - val_loss: 0.4909
Epoch 2/10
60/60          35s 534ms/step -
accuracy: 0.9119 - loss: 0.2598 - val_accuracy: 0.8892 - val_loss: 0.3831
Epoch 3/10
60/60          41s 541ms/step -
accuracy: 0.9332 - loss: 0.2067 - val_accuracy: 0.8834 - val_loss: 0.4602
Epoch 4/10
60/60          36s 542ms/step -
accuracy: 0.9444 - loss: 0.1809 - val_accuracy: 0.8761 - val_loss: 0.4500
Epoch 5/10
60/60          42s 560ms/step -
accuracy: 0.9521 - loss: 0.1361 - val_accuracy: 0.8717 - val_loss: 0.4385
Epoch 6/10
60/60          37s 572ms/step -


```

```
accuracy: 0.9581 - loss: 0.1139 - val_accuracy: 0.8732 - val_loss: 0.5043
Epoch 7/10
60/60          40s 546ms/step -
accuracy: 0.9586 - loss: 0.1278 - val_accuracy: 0.8732 - val_loss: 0.5187
Epoch 8/10
60/60          37s 566ms/step -
accuracy: 0.9682 - loss: 0.1029 - val_accuracy: 0.8673 - val_loss: 0.5314
Epoch 9/10
60/60          35s 538ms/step -
accuracy: 0.9698 - loss: 0.0847 - val_accuracy: 0.8673 - val_loss: 0.5903
Epoch 10/10
60/60          41s 537ms/step -
accuracy: 0.9820 - loss: 0.0596 - val_accuracy: 0.8878 - val_loss: 0.5326
```

```
[ ]: import matplotlib.pyplot as plt

# Plot training and validation accuracy
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.plot(history_fine_tune.history['accuracy'], label='Fine-tune Training_
↪Accuracy')
plt.plot(history_fine_tune.history['val_accuracy'], label='Fine-tune Validation_
↪Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')
plt.show()
```





```
[ ]: # Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(test_generator, steps=test_generator.
    ↪samples // batch_size)
print(f"Test Loss: {test_loss}")
print(f"Test Accuracy: {test_accuracy}")
```

```
4/4          1s 233ms/step -
accuracy: 0.8521 - loss: 0.5830
Test Loss: 0.7566991448402405
Test Accuracy: 0.8125
```

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
from tensorflow.keras.models import load_model

# Generate predictions for the test data without specifying 'steps'
test_predictions = model.predict(test_generator, verbose=1)

# Get the predicted class labels
```

```

predicted_labels = np.argmax(test_predictions, axis=1)

# Get the true class labels from the test generator
true_labels = test_generator.classes

# Generate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(10, 7))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=test_generator.
    class_indices.keys(), yticklabels=test_generator.class_indices.keys())
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()


```

WARNING:tensorflow:5 out of the last 11 calls to <function TensorFlowTrainer.make\_predict\_function.<locals>.one\_step\_on\_data\_distributed at 0x7e0cec2b75b0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce\_retracing=True option that can avoid unnecessary retracing. For (3), please refer to [https://www.tensorflow.org/guide/function#controlling\\_retracing](https://www.tensorflow.org/guide/function#controlling_retracing) and [https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) for more details.

4/5                      0s 113ms/step

WARNING:tensorflow:5 out of the last 11 calls to <function TensorFlowTrainer.make\_predict\_function.<locals>.one\_step\_on\_data\_distributed at 0x7e0cec2b75b0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce\_retracing=True option that can avoid unnecessary retracing. For (3), please refer to [https://www.tensorflow.org/guide/function#controlling\\_retracing](https://www.tensorflow.org/guide/function#controlling_retracing) and [https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) for more details.

5/5                      6s 785ms/step



```
[ ]: model.save("/content/drive/MyDrive/Flowers_Divided/MobileNetV2_fine_tune.keras")
```

### 3.3 Pre-trained ResNet50

```
[ ]: train_dir = '/content/drive/MyDrive/Flowers_Divided/train' # Path to your
      ↪train data
      val_dir = '/content/drive/MyDrive/Flowers_Divided/val' # Path to your
      ↪validation data
      test_dir = '/content/drive/MyDrive/Flowers_Divided/test'
```

```
[ ]: from tensorflow.keras.applications import ResNet50
      from tensorflow.keras.preprocessing.image import ImageDataGenerator
      from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout
      from tensorflow.keras.models import Sequential
      from tensorflow.keras.optimizers import Adam
      from tensorflow.keras.applications.resnet50 import preprocess_input as
      ↪resnet_preprocess
      # Data generators
      train_datagen = ImageDataGenerator(
```

```

        preprocessing_function=resnet_preprocess,
        rotation_range=20,
        width_shift_range=0.2,
        height_shift_range=0.2,
        shear_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True,
        fill_mode='nearest'
    )

    train_generator = train_datagen.flow_from_directory(
        train_dir,
        target_size=(224, 224),
        batch_size=32,
        class_mode='categorical'
    )

    val_generator = train_datagen.flow_from_directory(
        val_dir,
        target_size=(224, 224),
        batch_size=32,
        class_mode='categorical'
    )

```

Found 1919 images belonging to 5 classes.

Found 686 images belonging to 5 classes.

```

[ ]: # Model
base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
base_model.trainable = False

model = Sequential([
    base_model,
    GlobalAveragePooling2D(),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(5, activation='softmax')
])

model.compile(optimizer=Adam(learning_rate=1e-4),
              loss='categorical_crossentropy', metrics=['accuracy'])

model.summary()

```

Model: "sequential\_2"

Layer (type)	Output Shape	
↳Param #		
resnet50 (Functional)	(None, 7, 7, 2048)	
↳23,587,712		
global_average_pooling2d_2	(None, 2048)	
↳ 0		
(GlobalAveragePooling2D)		
↳		
dense_4 (Dense)	(None, 256)	
↳524,544		
dropout_2 (Dropout)	(None, 256)	
↳ 0		
dense_5 (Dense)	(None, 5)	
↳1,285		

Total params: 24,113,541 (91.99 MB)

Trainable params: 525,829 (2.01 MB)

Non-trainable params: 23,587,712 (89.98 MB)

### 3.3.1 Train

```
[ ]: # Train
import tensorflow as tf
from tensorflow.keras.callbacks import EarlyStopping, ReduceLRonPlateau,
↳ModelCheckpoint

LEARNING_RATE = 1e-4
early_stopping = EarlyStopping(monitor='val_loss', patience=5,
↳restore_best_weights=True)
reduce_lr = tf.keras.callbacks.ReduceLRonPlateau(monitor='val_loss', factor=0.
↳2, patience=3, min_lr=1e-6)
checkpoint = ModelCheckpoint(
    '/content/drive/MyDrive/flower_data/ResNet50.keras',          # Filepath
↳to save the model
    monitor='val_accuracy',      # Metric to monitor (e.g., validation accuracy)
    save_best_only=True,         # Save only the best model
```

```

    mode='max',                      # Maximize the monitored metric (e.g.,  $\uparrow$ 
    ↪accuracy)
    verbose=1                        # Print a message when saving the model
)
history = model.fit(train_generator, validation_data=val_generator,  $\uparrow$ 
    ↪epochs=25, callbacks=[early_stopping, reduce_lr, checkpoint])

```

Epoch 1/25

60/60 0s 476ms/step -

accuracy: 0.8191 - loss: 0.5075

Epoch 1: val\_accuracy improved from -inf to 0.86880, saving model to  
/content/drive/MyDrive/flower\_data/ResNet50.keras

60/60 45s 702ms/step -

accuracy: 0.8193 - loss: 0.5071 - val\_accuracy: 0.8688 - val\_loss: 0.4113 -  
learning\_rate: 1.0000e-04

Epoch 2/25

60/60 0s 491ms/step -

accuracy: 0.8240 - loss: 0.4618

Epoch 2: val\_accuracy did not improve from 0.86880

60/60 81s 684ms/step -

accuracy: 0.8242 - loss: 0.4613 - val\_accuracy: 0.8659 - val\_loss: 0.3927 -  
learning\_rate: 1.0000e-04

Epoch 3/25

60/60 0s 481ms/step -

accuracy: 0.8673 - loss: 0.3743

Epoch 3: val\_accuracy did not improve from 0.86880

60/60 83s 696ms/step -

accuracy: 0.8673 - loss: 0.3744 - val\_accuracy: 0.8601 - val\_loss: 0.3866 -  
learning\_rate: 1.0000e-04

Epoch 4/25

60/60 0s 481ms/step -

accuracy: 0.8929 - loss: 0.3320

Epoch 4: val\_accuracy improved from 0.86880 to 0.88192, saving model to  
/content/drive/MyDrive/flower\_data/ResNet50.keras

60/60 55s 844ms/step -

accuracy: 0.8928 - loss: 0.3324 - val\_accuracy: 0.8819 - val\_loss: 0.3532 -  
learning\_rate: 1.0000e-04

Epoch 5/25

60/60 0s 480ms/step -

accuracy: 0.8795 - loss: 0.3154

Epoch 5: val\_accuracy did not improve from 0.88192

60/60 71s 675ms/step -

accuracy: 0.8795 - loss: 0.3157 - val\_accuracy: 0.8732 - val\_loss: 0.3577 -  
learning\_rate: 1.0000e-04

Epoch 6/25

60/60 0s 469ms/step -

accuracy: 0.8858 - loss: 0.3238

Epoch 6: val\_accuracy did not improve from 0.88192

```

60/60          43s 663ms/step -
accuracy: 0.8860 - loss: 0.3233 - val_accuracy: 0.8790 - val_loss: 0.3439 -
learning_rate: 1.0000e-04
Epoch 7/25
60/60          0s 470ms/step -
accuracy: 0.8855 - loss: 0.3061
Epoch 7: val_accuracy improved from 0.88192 to 0.89504, saving model to
/content/drive/MyDrive/flower_data/ResNet50.keras
60/60          44s 681ms/step -
accuracy: 0.8856 - loss: 0.3058 - val_accuracy: 0.8950 - val_loss: 0.3212 -
learning_rate: 1.0000e-04
Epoch 8/25
60/60          0s 482ms/step -
accuracy: 0.9042 - loss: 0.2672
Epoch 8: val_accuracy did not improve from 0.89504
60/60          44s 674ms/step -
accuracy: 0.9042 - loss: 0.2673 - val_accuracy: 0.8834 - val_loss: 0.3257 -
learning_rate: 1.0000e-04
Epoch 9/25
60/60          0s 470ms/step -
accuracy: 0.9109 - loss: 0.2470
Epoch 9: val_accuracy did not improve from 0.89504
60/60          91s 817ms/step -
accuracy: 0.9110 - loss: 0.2470 - val_accuracy: 0.8776 - val_loss: 0.3272 -
learning_rate: 1.0000e-04
Epoch 10/25
60/60          0s 472ms/step -
accuracy: 0.9217 - loss: 0.2233
Epoch 10: val_accuracy improved from 0.89504 to 0.89942, saving model to
/content/drive/MyDrive/flower_data/ResNet50.keras
60/60          73s 684ms/step -
accuracy: 0.9216 - loss: 0.2235 - val_accuracy: 0.8994 - val_loss: 0.3191 -
learning_rate: 1.0000e-04
Epoch 11/25
60/60          0s 479ms/step -
accuracy: 0.9143 - loss: 0.2328
Epoch 11: val_accuracy did not improve from 0.89942
60/60          44s 673ms/step -
accuracy: 0.9144 - loss: 0.2326 - val_accuracy: 0.8878 - val_loss: 0.3172 -
learning_rate: 1.0000e-04
Epoch 12/25
60/60          0s 466ms/step -
accuracy: 0.9209 - loss: 0.2190
Epoch 12: val_accuracy did not improve from 0.89942
60/60          42s 656ms/step -
accuracy: 0.9209 - loss: 0.2190 - val_accuracy: 0.8848 - val_loss: 0.3188 -
learning_rate: 1.0000e-04
Epoch 13/25

```

```

60/60          0s 493ms/step -
accuracy: 0.9367 - loss: 0.1847
Epoch 13: val_accuracy improved from 0.89942 to 0.90525, saving model to
/content/drive/MyDrive/flower_data/ResNet50.keras
60/60          45s 701ms/step -
accuracy: 0.9367 - loss: 0.1848 - val_accuracy: 0.9052 - val_loss: 0.2887 -
learning_rate: 1.0000e-04
Epoch 14/25
60/60          0s 484ms/step -
accuracy: 0.9361 - loss: 0.1874
Epoch 14: val_accuracy did not improve from 0.90525
60/60          54s 833ms/step -
accuracy: 0.9361 - loss: 0.1876 - val_accuracy: 0.8950 - val_loss: 0.2860 -
learning_rate: 1.0000e-04
Epoch 15/25
60/60          0s 478ms/step -
accuracy: 0.9294 - loss: 0.1967
Epoch 15: val_accuracy did not improve from 0.90525
60/60          43s 653ms/step -
accuracy: 0.9294 - loss: 0.1966 - val_accuracy: 0.8907 - val_loss: 0.2909 -
learning_rate: 1.0000e-04
Epoch 16/25
60/60          0s 472ms/step -
accuracy: 0.9516 - loss: 0.1527
Epoch 16: val_accuracy did not improve from 0.90525
60/60          81s 663ms/step -
accuracy: 0.9515 - loss: 0.1529 - val_accuracy: 0.9023 - val_loss: 0.2874 -
learning_rate: 1.0000e-04
Epoch 17/25
60/60          0s 472ms/step -
accuracy: 0.9512 - loss: 0.1589
Epoch 17: val_accuracy did not improve from 0.90525
60/60          83s 681ms/step -
accuracy: 0.9512 - loss: 0.1588 - val_accuracy: 0.9038 - val_loss: 0.2810 -
learning_rate: 1.0000e-04
Epoch 18/25
60/60          0s 470ms/step -
accuracy: 0.9476 - loss: 0.1537
Epoch 18: val_accuracy did not improve from 0.90525
60/60          43s 660ms/step -
accuracy: 0.9476 - loss: 0.1538 - val_accuracy: 0.8878 - val_loss: 0.2996 -
learning_rate: 1.0000e-04
Epoch 19/25
60/60          0s 468ms/step -
accuracy: 0.9405 - loss: 0.1688
Epoch 19: val_accuracy did not improve from 0.90525
60/60          92s 816ms/step -
accuracy: 0.9405 - loss: 0.1688 - val_accuracy: 0.8907 - val_loss: 0.3142 -

```



```

learning_rate: 1.0000e-04
Epoch 20/25
60/60          0s 466ms/step -
accuracy: 0.9485 - loss: 0.1590
Epoch 20: val_accuracy did not improve from 0.90525
60/60          71s 658ms/step -
accuracy: 0.9484 - loss: 0.1590 - val_accuracy: 0.8892 - val_loss: 0.3077 -
learning_rate: 1.0000e-04
Epoch 21/25
60/60          0s 495ms/step -
accuracy: 0.9503 - loss: 0.1477
Epoch 21: val_accuracy did not improve from 0.90525
60/60          83s 683ms/step -
accuracy: 0.9503 - loss: 0.1477 - val_accuracy: 0.8950 - val_loss: 0.2890 -
learning_rate: 2.0000e-05
Epoch 22/25
60/60          0s 470ms/step -
accuracy: 0.9484 - loss: 0.1452
Epoch 22: val_accuracy did not improve from 0.90525
60/60          81s 654ms/step -
accuracy: 0.9484 - loss: 0.1452 - val_accuracy: 0.8994 - val_loss: 0.2899 -
learning_rate: 2.0000e-05

```

```

[ ]: test_datagen = ImageDataGenerator(preprocessing_function=resnet_preprocess) #
      ↪ Replace with your model's preprocessing

test_generator = test_datagen.flow_from_directory(
    test_dir, # Path to your test data
    target_size=(224, 224), # Same as training
    batch_size=32, # Adjust as needed
    class_mode='categorical', # Same as training
    shuffle=False # Do not shuffle for evaluation
)

```

Found 141 images belonging to 5 classes.

```

[ ]: # Evaluate the model
results = model.evaluate(test_generator)
print(f"Test Loss: {results[0]}")
print(f"Test Accuracy: {results[1]}")

```

```

/usr/local/lib/python3.10/dist-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
    self._warn_if_super_not_called()

```

5/5                    1s 186ms/step -  
accuracy: 0.9525 - loss: 0.1839  
Test Loss: 0.264686644077301  
Test Accuracy: 0.9290780425071716

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
from tensorflow.keras.models import load_model

# Generate predictions for the test data without specifying 'steps'
test_predictions = model.predict(test_generator, verbose=1)


# Get the predicted class labels
predicted_labels = np.argmax(test_predictions, axis=1)

# Get the true class labels from the test generator
true_labels = test_generator.classes

# Generate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(10, 7))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=test_generator.
    class_indices.keys(), yticklabels=test_generator.class_indices.keys())
plt.title("Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```

5/5                    2s 121ms/step



## 4 Comparison of Pre-trained Models

**Detailed Analysis 1.** Complexity MobileNetV2 and EfficientNet-B0 are lightweight, making them suitable for resource-constrained devices like mobile phones or embedded systems.

ResNet-50 and InceptionV3 offer a balance between complexity and performance.

VGG-16 and DenseNet-121 are heavier models with higher memory and computational requirements.

Vision Transformers (ViT) are complex and require substantial resources for training but excel in handling large datasets and structured data.

2. Parameters VGG-16 has the highest number of parameters, making it memory-intensive and slow during inference.

MobileNetV2 and NASNetMobile are highly efficient with fewer parameters, making them ideal for edge devices.

ResNet-50 and InceptionV3 strike a good balance, providing robust performance without being excessively large.

3. Key Features MobileNetV2: Depthwise separable convolutions drastically reduce computation without sacrificing performance.

EfficientNet-B0: Uses a compound scaling method to scale depth, width, and resolution effectively.

ResNet-50: Residual connections alleviate the vanishing gradient problem, enabling deep architectures.

DenseNet-121: Reuses features, reducing redundancy and improving gradient flow.

InceptionV3: Factorized convolutions and auxiliary classifiers enhance learning.

4. Problems to Solve Simple Classification Tasks: MobileNetV2, EfficientNet-B0, and NASNet-Mobile are ideal for classification tasks on small or medium datasets.

High-performance Applications: ResNet, Inception, and DenseNet perform well for larger and more complex datasets.

Sequential/Structured Data: Vision Transformers (ViT) excel in structured or sequential image data processing, like video analytics or multi-task learning.

5. Use Cases Mobile and Embedded Devices: MobileNetV2, NASNetMobile, and EfficientNet-B0 are optimized for low-power environments.

Feature Extraction: ResNet-50 and DenseNet-121 are commonly used as backbone models for feature extraction in custom models.

Research and Development: VGG-16 and Vision Transformers are favored in academic and experimental scenarios.

Object Detection and Segmentation: InceptionV3 and DenseNet-121 are often used for advanced tasks like object detection and semantic segmentation.

When to Choose Which Model? Resource Constraints: Use MobileNetV2 or EfficientNet-B0 for devices with limited computation or memory.

Large Datasets: Use ResNet-50, InceptionV3, or Vision Transformers for tasks requiring high accuracy on complex datasets.

Real-time Applications: Use MobileNetV2 or NASNetMobile for applications like real-time image recognition.

Advanced Research: Use Vision Transformers or DenseNet-121 for innovative applications in structured data processing or semantic segmentation.

[ ]: